

# Machine-Level Programming: Advanced Topics

CSCI3240: Lecture 11

Dr. Arpan Man Sainju

Middle Tennessee State University

# Today

- **Memory Layout**
- **Buffer Overflow**
  - Vulnerability
  - Protection

# x86-64 Linux Memory Layout *not drawn to scale*

00007FFFFFFF

## ■ Stack

- Runtime stack (8MB limit)
- E. g., local variables

## ■ Heap


- Dynamically allocated as needed
- When call `malloc()`, `calloc()`, `new()`

## ■ Data

- Statically allocated data
- E.g., global vars, `static` vars, string constants

## ■ Text / Shared Libraries

- Executable machine instructions
- Read-only



# Today

- Memory Layout
- **Buffer Overflow**
  - Vulnerability
  - Protection

# Recall: Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

```
fun(0)    ↪ 3.14
fun(1)    ↪ 3.14
fun(2)    ↪ 3.13999998664856
fun(3)    ↪ 2.000000061035156
fun(4)    ↪ 3.14
fun(6)    ↪ Segmentation fault
```

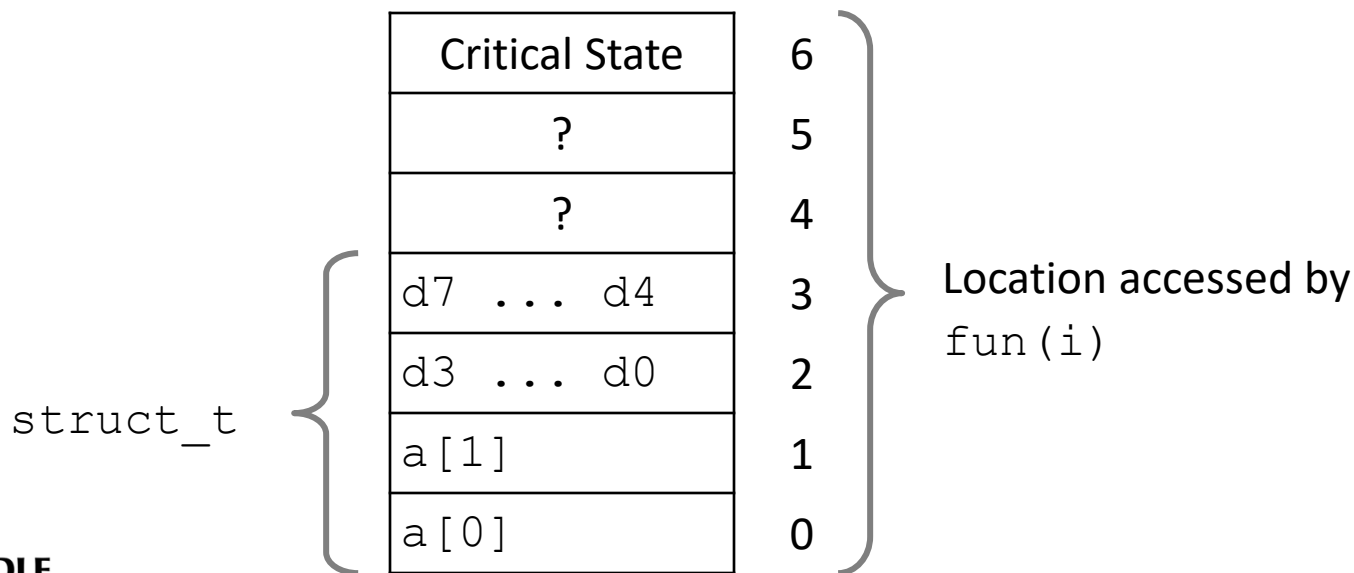
- Result is system specific

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)    ↪ 3.14  
fun(1)    ↪ 3.14  
fun(2)    ↪ 3.1399998664856  
fun(3)    ↪ 2.00000061035156  
fun(4)    ↪ 3.14  
fun(6)    ↪ Segmentation fault

## Explanation:



# Such problems are a BIG deal

- **Generally called a “buffer overflow”**

- when exceeding the memory size allocated for an array

- **Why a big deal?**

- It's the #1 technical cause of security vulnerabilities
  - #1 overall cause is social engineering / user ignorance

- **Most common form**

- Unchecked lengths on string inputs
- Particularly for bounded character arrays on the stack
  - sometimes referred to as stack smashing

# String Library Code

## ■ Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read

## ■ Similar problems with other library functions

- **strcpy, strcat**: Copy strings of arbitrary length
- **scanf, fscanf, sscanf**, when given **%s** conversion specification



# Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

← btw, how big  
is big enough?

```
void call_echo() {  
    echo();  
}
```

```
unix>./bufdemo-nsp  
Type a string:012345678901234567890123  
012345678901234567890123
```

```
unix>./bufdemo-nsp  
Type a string:0123456789012345678901234  
Segmentation Fault
```

# Buffer Overflow Disassembly

echo:

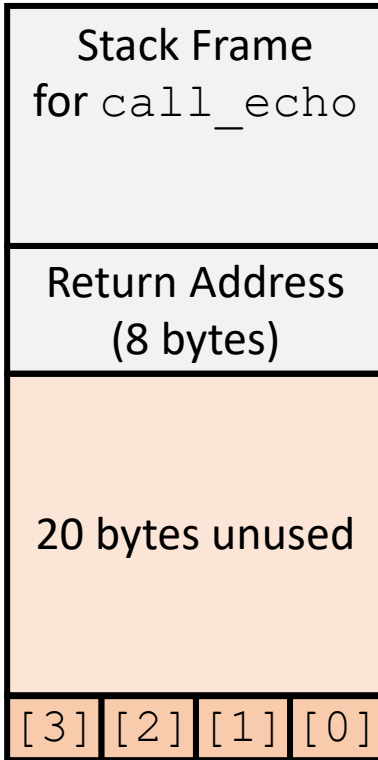
```
00000000004006cf <echo>:
 4006cf:  48 83 ec 18          sub     $0x18,%rsp
 4006d3:  48 89 e7            mov     %rsp,%rdi
 4006d6:  e8 a5 ff ff ff     callq  400680 <gets>
 4006db:  48 89 e7            mov     %rsp,%rdi
 4006de:  e8 3d fe ff ff     callq  400520 <puts@plt>
 4006e3:  48 83 c4 18          add     $0x18,%rsp
 4006e7:  c3                  retq
```

call\_echo:

```
4006e8:  48 83 ec 08          sub     $0x8,%rsp
4006ec:  b8 00 00 00 00      mov     $0x0,%eax
4006f1:  e8 d9 ff ff ff     callq  4006cf <echo>
4006f6:  48 83 c4 08          add     $0x8,%rsp
4006fa:  c3                  retq
```

# Buffer Overflow Stack

*Before call to gets*




```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

# Buffer Overflow Stack Example

*Before call to gets*



```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

call\_echo:

```
. . .  
4006f1:    callq    4006cf <echo>  
4006f6:    add      $0x8, %rsp  
. . .
```

# Buffer Overflow Stack Example #1

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()
{
    char buf[4];
    gets(buf);
    . . .
}
```

```
echo:
    subq    $24, %rsp
    movq    %rsp, %rdi
    call    gets
    . . .
```

call\_echo:

```
. . .
4006f1:    callq   4006cf <echo>
4006f6:    add     $0x8,%rsp
. . .
```

buf ← %rsp

```
unix> ./bufdemo-nsp
Type a string: 01234567890123456789012
01234567890123456789012
```

Overflowed buffer, but did not corrupt state

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

# Buffer Overflow Stack Example #2

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

call\_echo:

```
. . .  
4006f1:    callq   4006cf <echo>  
4006f6:    add     $0x8,%rsp  
. . .
```

buf ← %rsp

```
unix> ./bufdemo-nsp  
Type a string: 0123456789012345678901234  
Segmentation Fault
```

Overflowed buffer and corrupted return pointer

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

# Buffer Overflow Stack Example #3

After call to gets

Stack Frame for call_echo			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()
{
    char buf[4];
    gets(buf);
    . . .
}
```

```
echo:
    subq    $24, %rsp
    movq    %rsp, %rdi
    call    gets
    . . .
```

call\_echo:

```
. . .
4006f1:    callq    4006cf <echo>
4006f6:    add      $0x8,%rsp
. . .
```

buf ← %rsp

```
unix> ./bufdemo-nsp
Type a string: 012345678901234567890123
012345678901234567890123
```

Overflowed buffer, corrupted return pointer, but program seems to work!

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

# Buffer Overflow Stack Example #3 Explained

*After call to gets*

Stack Frame for call_echo			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

buf ← %rsp

register\_tm\_clones:

```
. . .  
400600:  mov    %rsp,%rbp  
400603:  mov    %rax,%rdx  
400606:  shr    $0x3f,%rdx  
40060a:  add    %rdx,%rax  
40060d:  sar    %rax  
400610:  jne    400614  
400612:  pop    %rbp  
400613:  retq
```

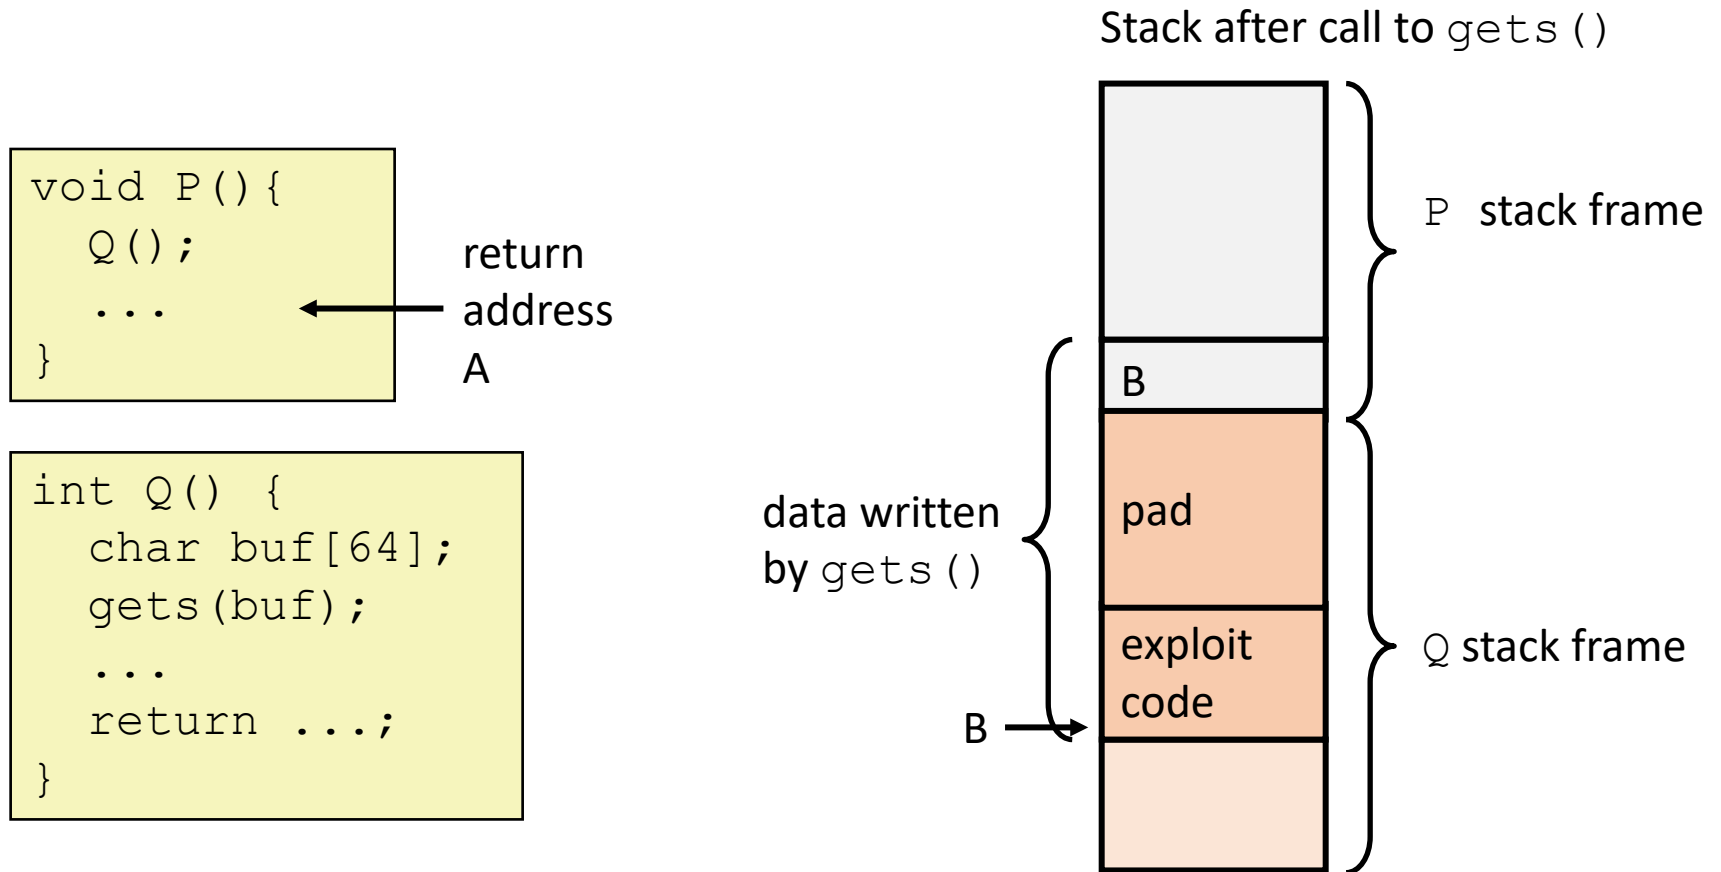
“Returns” to unrelated code

Lots of things happen, without modifying critical state

Eventually executes `retq` back to `main`



# Code Injection Attacks



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code

# Exploits Based on Buffer Overflows

- *Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines*
  
- **Distressingly common in real programs**
  - Programmers keep making the same mistakes ☹
  - Recent measures make these attacks much more difficult
  
- **Examples across the decades**
  - Original “Internet worm” (1988)
  - “IM wars” (1999)
  - Twilight hack on Wii (2000s)
  - ... and many, many more

# Example: the original Internet worm (1988)

## ▪ Exploited a few vulnerabilities to spread

- Early versions of the finger server used `gets()` to read the argument sent by the client:
  - `finger droh@cs.cmu.edu`
- Worm attacked finger server by sending phony argument:
  - `finger "exploit-code padding new-return-address"`
  - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

## ▪ Once on a machine, scanned for other machines to attack

- invaded ~6000 computers in hours (10% of the Internet 😊 )
  - see June 1989 article in *Comm. of the ACM*
- the young author of the worm was prosecuted...
- and CERT (Computer Emergency Response Team) was formed

# Aside: Worms and Viruses

- **Worm: A program that**
  - Can run by itself
  - Can propagate a fully working version of itself to other computers
  
- **Virus: Code that**
  - Adds itself to other programs
  - Does not run independently
  
- **Both are (usually) designed to spread among computers and to wreak havoc**

# OK, what to do about buffer overflow attacks

- Avoid overflow vulnerabilities
- Employ system-level protections
- Have compiler use “stack canaries”
- Lets talk about each...

# 1. Avoid Overflow Vulnerabilities in Code (!)

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```


- For example, use library routines that limit string lengths
  - **fgets** instead of **gets**
  - **strncpy** instead of **strcpy**
  - Don't use **scanf** with **%s** conversion specification
    - Use **fgets** to read the string
    - Or use **%ns** where **n** is a suitable integer

## 2. System-Level Protections can help

### ▪ Randomized stack offsets

- At start of program, allocate random amount of space on stack
- Shifts stack addresses for entire program
- Makes it difficult for hacker to predict beginning of inserted code
- E.g.: 5 executions of memory allocation code
- Stack repositioned each time program executes


local      0x7ffe4d3be87c    0x7fff75a4f9fc    0x7ffeadb7c80c    0x7ffeaea2fdac    0x7ffcd452017c



## 2. System-Level Protections can help

### ▪ Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
  - Can execute anything readable
- X86-64 added explicit “execute” permission
- Stack marked as non-executable



Any attempt to execute this code will fail



# 3. Stack Canaries can help


## ■ Idea

- Place special value (“canary”) on stack just beyond buffer
- Check for corruption before exiting function

## ■ GCC Implementation

- **-fstack-protector**
- Now the default

Compiling with stack protector



```
unix> ./bufdemo-sp
Type a string: 0123456
0123456
```

```
unix> ./bufdemo-sp
Type a string: 01234567
*** stack smashing detected ***
```

*Bufdemo.c*

```
/* Demonstration of buffer overflow */

#include <stdio.h>
#include <stdlib.h>

/* Read input line and write it back */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}

void call_echo() {
    echo();
}

int main()
{
    printf("Type a string:");
    call_echo();
    return 0;
}
```

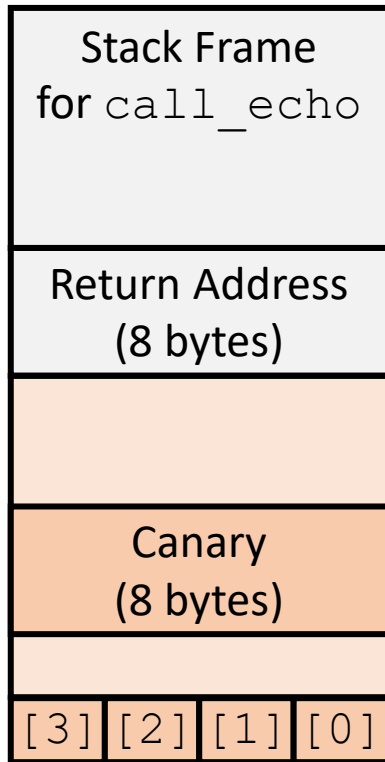
# Protected Buffer Disassembly

echo:

```
40072f:  sub    $0x18,%rsp
400733:  mov     %fs:0x28,%rax
40073c:  mov     %rax,0x8(%rsp)
400741:  xor     %eax,%eax
400743:  mov     %rsp,%rdi
400746:  callq   4006e0 <gets>
40074b:  mov     %rsp,%rdi
40074e:  callq   400570 <puts@plt>
400753:  mov     0x8(%rsp),%rax
400758:  xor     %fs:0x28,%rax
400761:  je      400768 <echo+0x39>
400763:  callq   400580 <__stack_chk_fail@plt>
400768:  add     $0x18,%rsp
40076c:  retq
```

# Setting Up Canary

*Before call to gets*



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    . . .
    movq    %fs:40, %rax    # Get canary
    movq    %rax, 8(%rsp)  # Place on stack
    xorl    %eax, %eax     # Erase canary
    . . .
```

# Checking Canary

After call to gets

Stack Frame for call_echo			
Return Address (8 bytes)			
Canary (8 bytes)			
00	36	35	34
33	32	31	30

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

Input: 0123456

buf ← %rsp

```
echo:
    . . .
    movq    8(%rsp), %rax    # Retrieve from
stack
    xorq    %fs:40, %rax    # Compare to canary
    je      .L6             # If same, OK
    call    __stack_chk_fail # FAIL
```

.L6:

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition