

Chapter 11: Network Programming


CSCI3240: Lecture 14 and 15

Dr. Arpan Man Sainju

Middle Tennessee State University

Basic Terminology

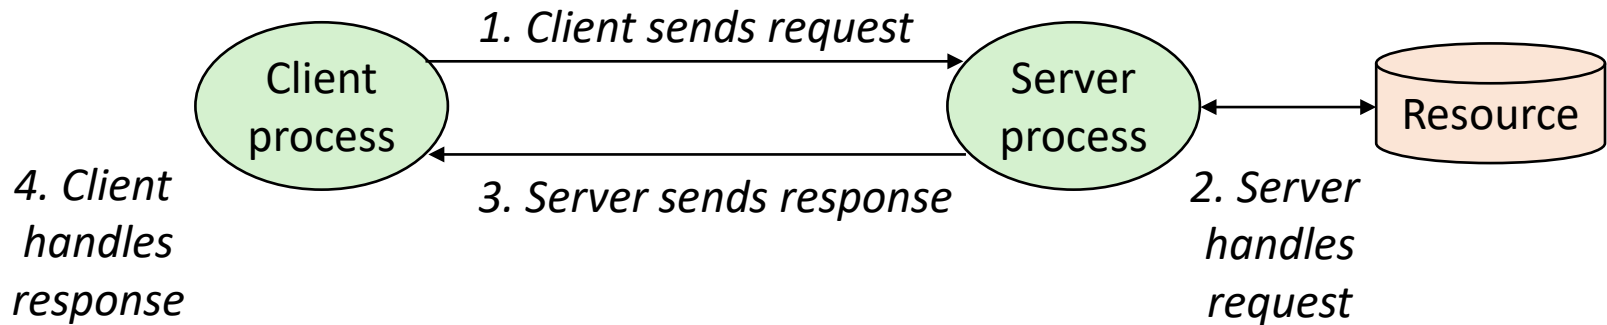
- Computer Network: A computer network is composed of multiple computers connected using a telecommunication system.
- Host: machines
 - PCs, workstation
 - Network components: routers
- Interconnection maybe any medium capable of communicating information:
 - Cooper wire (Ethernet)
 - Lasers (Optical fiber)
 - Radio/Satellite link
 - Cable (coax)



A Client-Server Transaction

- **Most network applications are based on the client-server model:**
 - A **server** process and one or more **client** processes
 - Server manages some **resource**
 - Server provides **service** by manipulating resource for clients
 - Server activated by request from client (vending machine analogy)
- **A server is a process – not a machine.**
- **A server waits for a request from a client.**
- **A client is a process that sends a request to an existing server and (usually) waits for a reply.**

A Client-Server Transaction



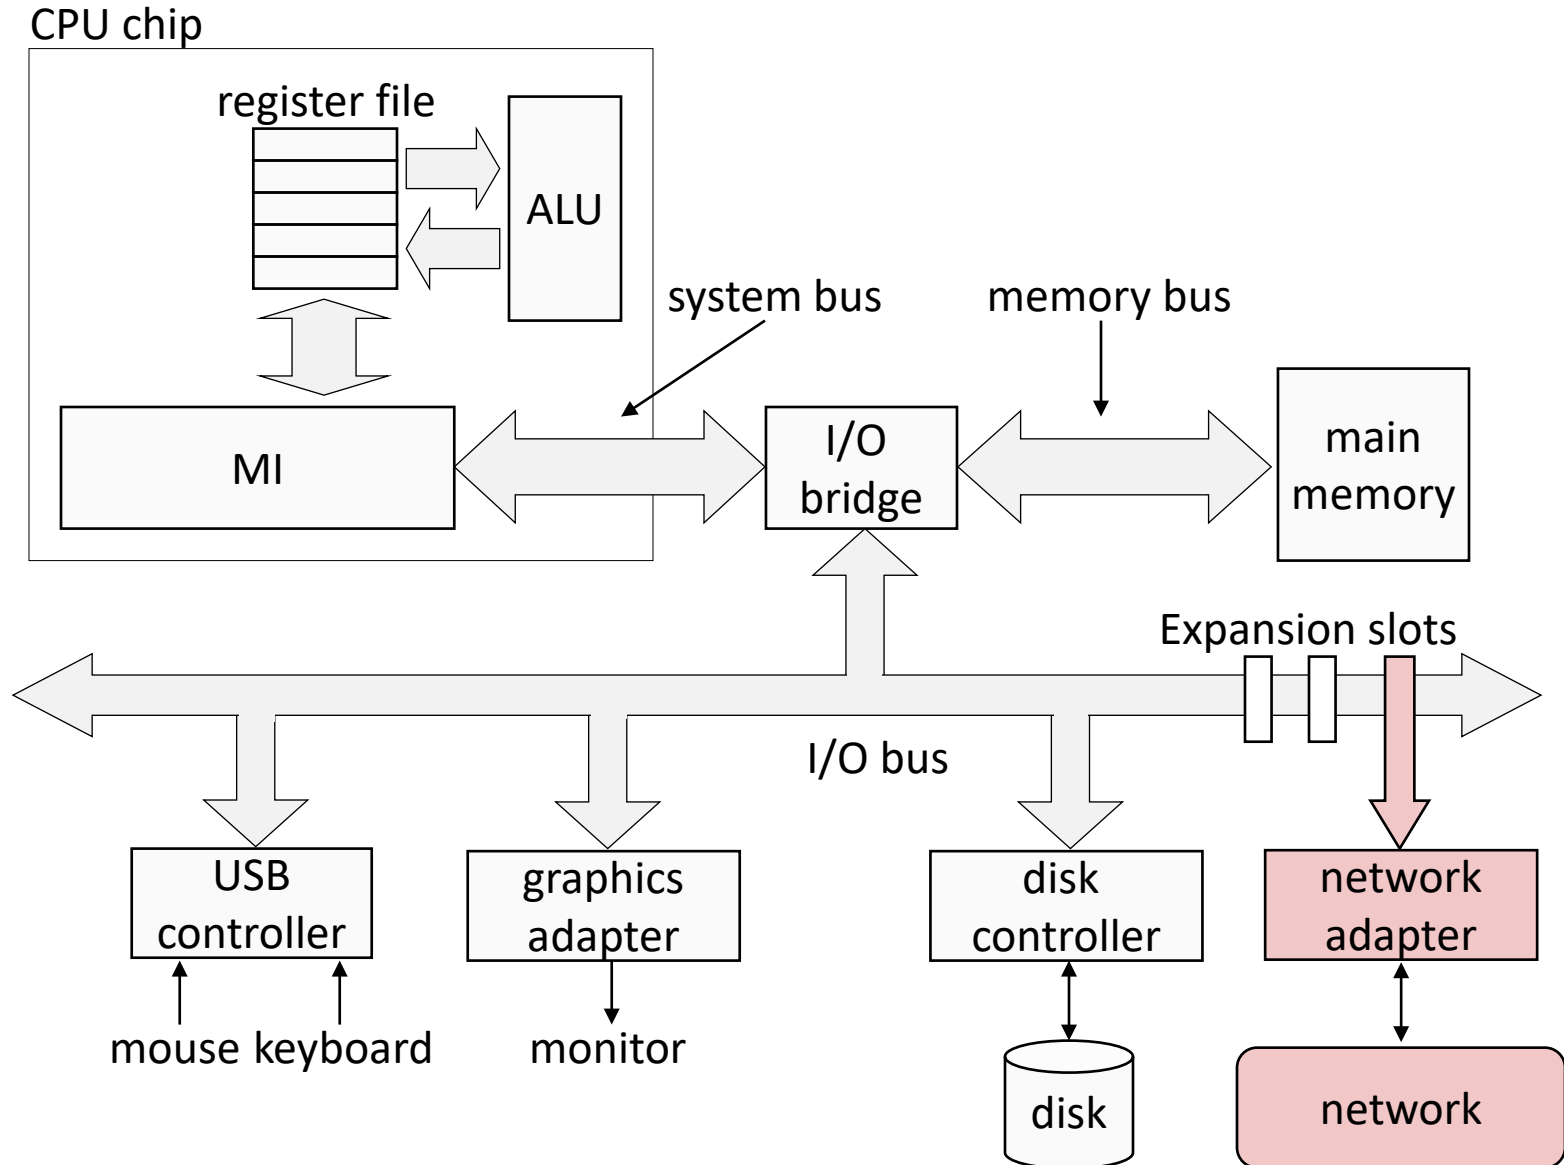
*Note: clients and servers are processes running on hosts
(can be the same or different hosts)*

Example: A Web server manages a set of disk files that it retrieves and executes on behalf of clients.

Client-Server Examples

- **Web Server:** Client request for a web page, server delivers
- **Mail Server:** Email servers can be used for sending and receiving emails.
- **File Servers:** They are the centralized locations for the files.
Example: One drive, Google drive.


Hardware Organization of a Network Host



Computer Networks

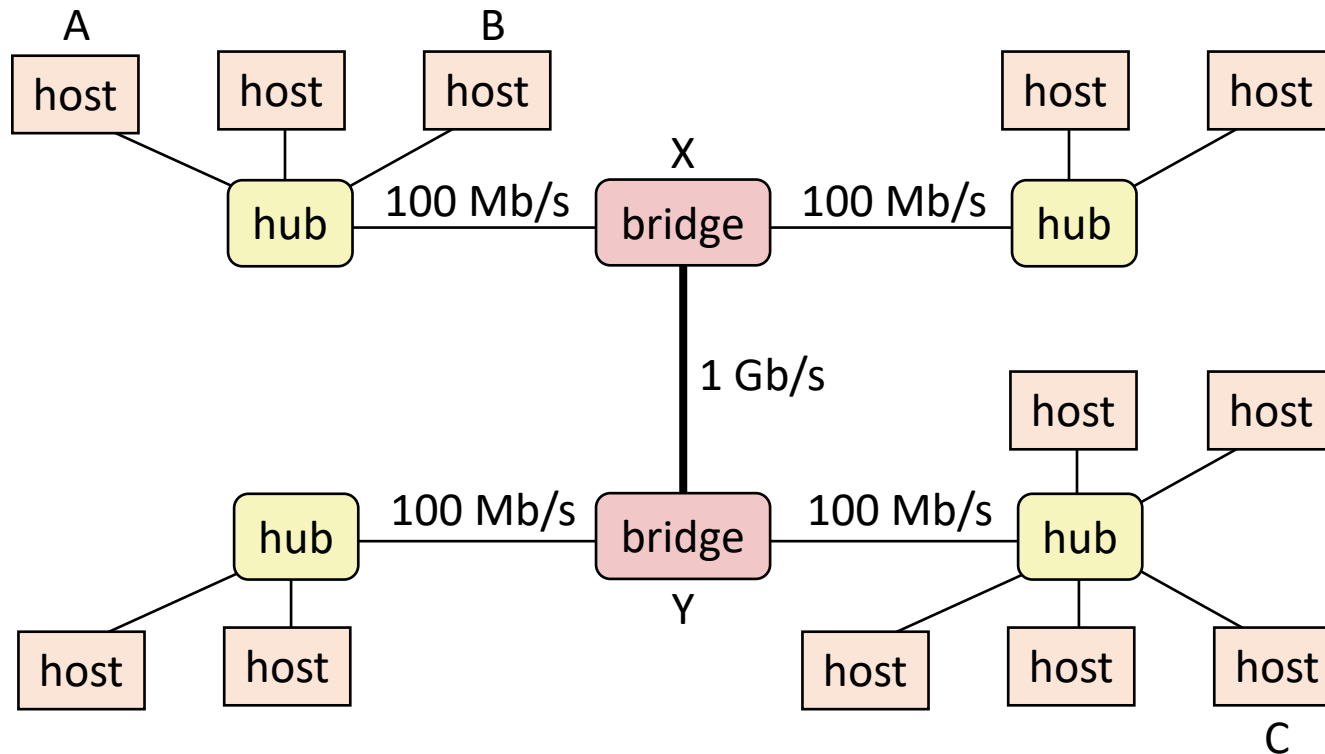
- A ***network*** is a hierarchical system of boxes and wires organized by geographical proximity
 - SAN (System Area Network) spans cluster or machine room
 - Switched Ethernet,
 - LAN (Local Area Network) spans a building or campus
 - Ethernet is most prominent example
 - WAN (Wide Area Network) spans country or world
 - Typically, high-speed point-to-point phone lines
- An ***internetwork (internet)*** is an interconnected set of networks
 - The Global IP Internet (uppercase “I”) is the most famous example of an internet (lowercase “i”)
- Let’s see how an internet is built from the ground up

Lowest Level: Ethernet Segment



- Ethernet segment consists of a collection of *hosts* connected by wires (twisted pairs) to a *hub*
- Spans room or floor in a building
- Operation
 - Each Ethernet adapter has a unique 48-bit address (MAC address)
 - E.g., 00:16:ea:e3:54:e6
 - Hosts send bits to any other host in chunks called *frames*
 - Hub copies each bit from each port to every other port
 - Every host sees every bit
 - Note: Hubs are on their way out. Bridges (switches, routers) became cheap enough to replace them


Next Level: Bridged Ethernet Segment



- Spans building or campus
- Bridges cleverly learn which hosts are reachable from which ports and then selectively copy frames from port to port


Conceptual View of LANs

- For simplicity, hubs, bridges, and wires are often shown as a collection of hosts attached to a single wire:




Next Level: internets

- Multiple incompatible LANs can be physically connected by specialized computers called *routers*
- The connected networks are called an *internet* (lower case)



LAN 1 and LAN 2 might be completely different, totally incompatible (e.g., Ethernet, Fiber Channel, 802.11, T1-links, DSL, ...)*

Logical Structure of an internet



- **Ad hoc interconnection of networks**

- No particular topology
- Vastly different router & link capacities

- **Send packets from source to destination by hopping through networks**


- Router forms bridge from one network to another
- Different packets may take different routes

The Notion of an internet Protocol

- How is it possible to send bits across incompatible LANs and WANs?
- Solution: *protocol* software running on each host and router
 - Protocol is a set of rules that governs how hosts and routers should cooperate when they transfer data from network to network.
 - Smooths out the differences between the different networks

What's a protocol?

- a human protocol and a computer network protocol:



protocols define format, order of msgs sent and received among network entities, and actions taken on msg transmission, receipt

What Does an internet Protocol Do?

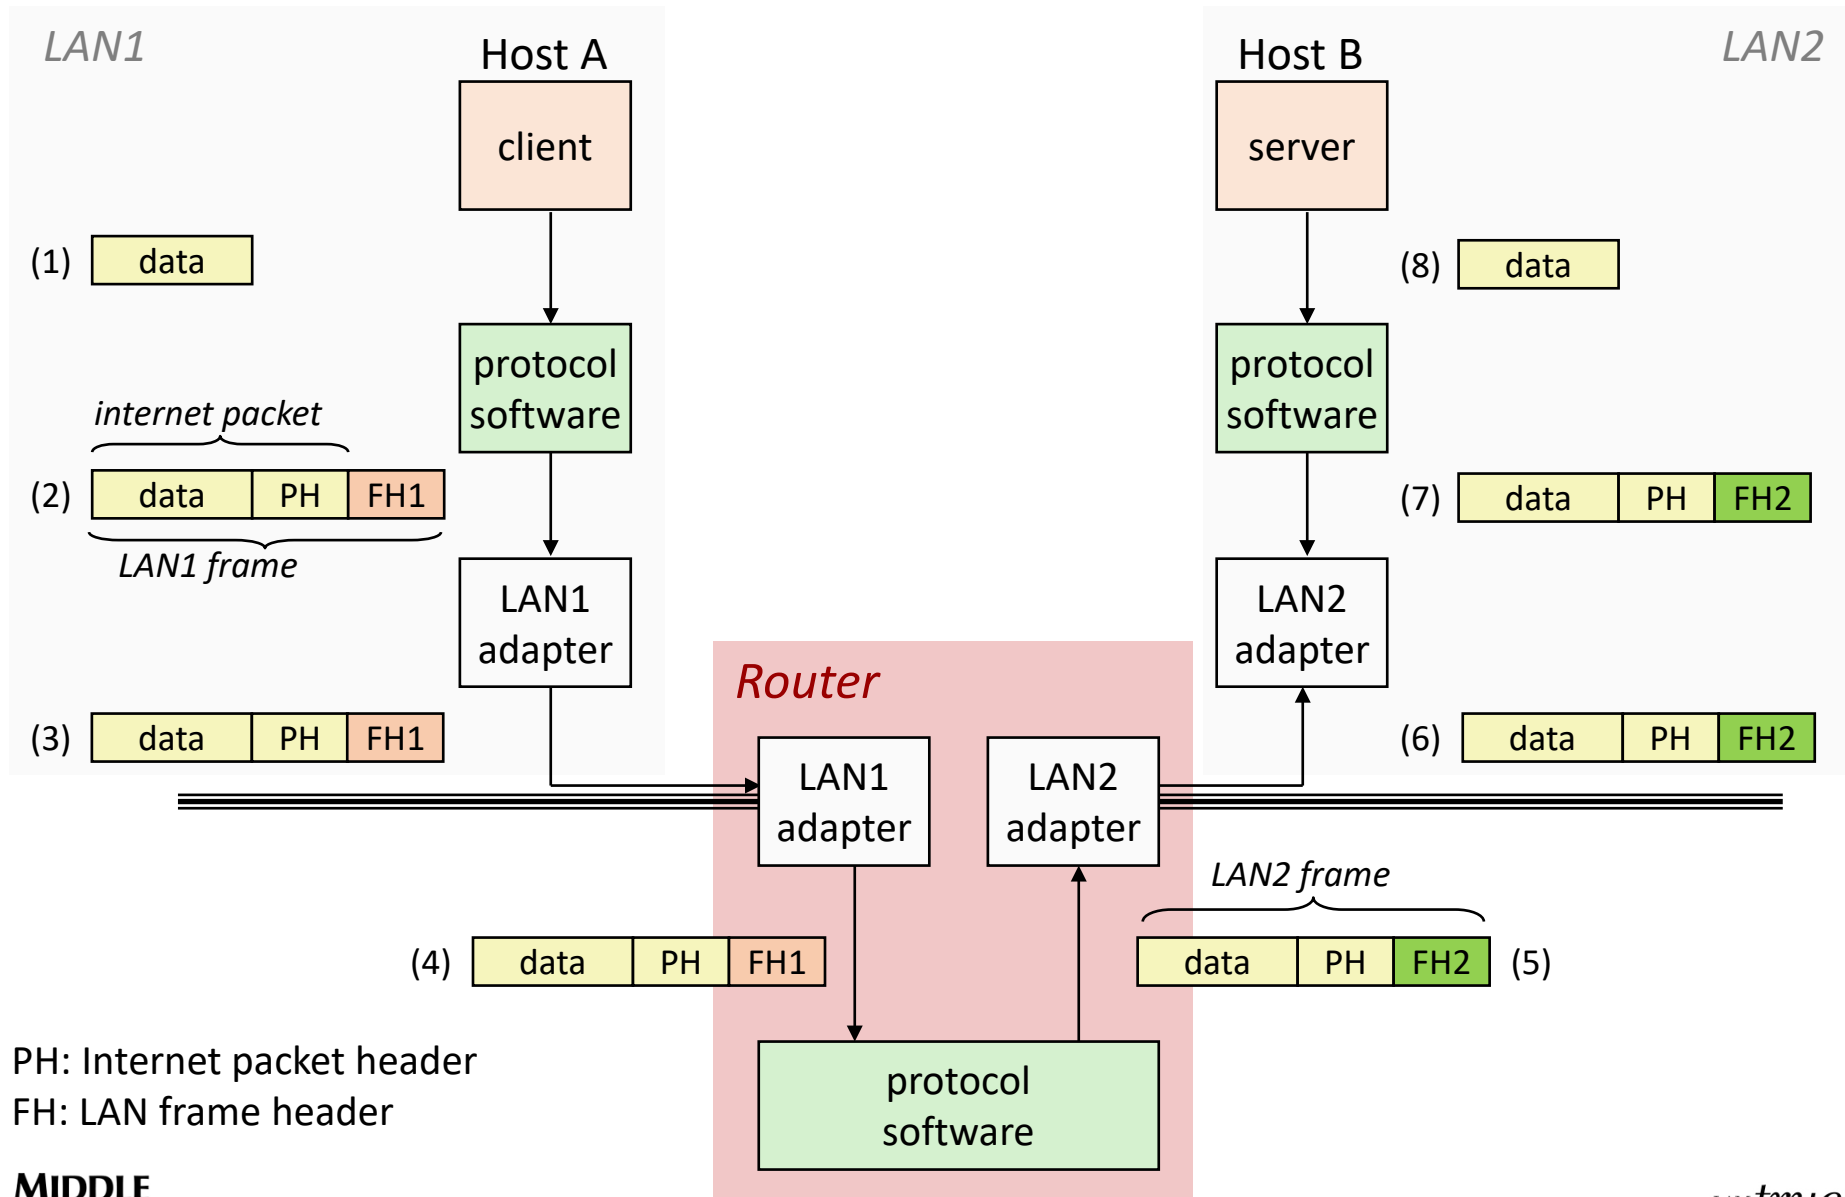
■ Provides a *naming scheme*

- Different LAN technologies have different and incompatible ways of assigning address to host.
- An internet protocol smooths these differences by defining a uniform format for *host addresses*
- Each host (and router) is assigned at least one of these internet addresses that uniquely identifies it

■ Provides a *delivery mechanism*

- An internet protocol defines a standard transfer unit (*packet*)
- Packet consists of *header* and *payload*
 - Header: contains info such as packet size, source, and destination addresses
 - Payload: contains data bits sent from the source host

Transferring internet Data Via Encapsulation



Other Issues

- **We are glossing over a number of important questions:**
 - What if different networks have different maximum frame sizes? (segmentation)
 - How do routers know where to forward frames?
 - How are routers informed when the network topology changes?
 - What if packets get lost?

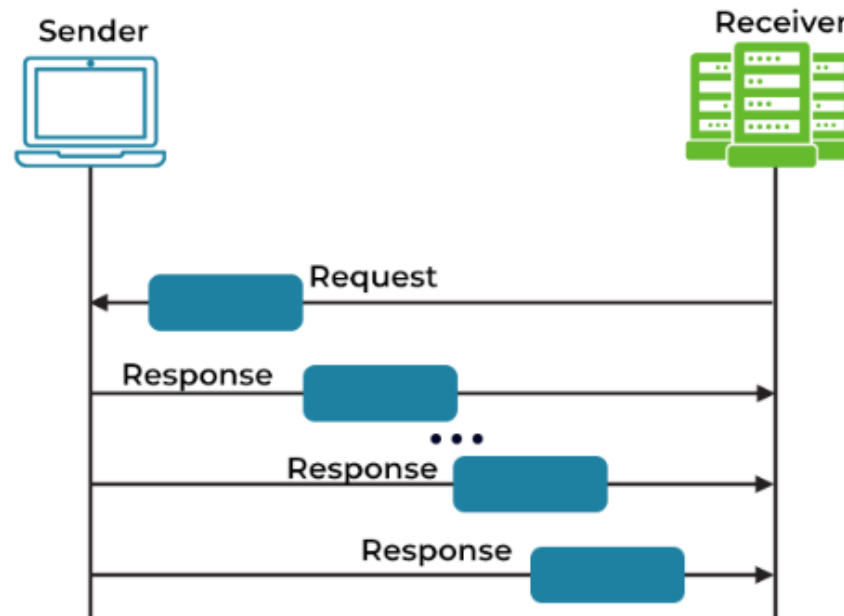
- **These (and other) questions are addressed by the area of systems known as *computer networking***
 - *CSCI 4300 – Data Communication and Network*

Global IP Internet (upper case)

- Most famous example of an internet
- Based on the TCP/IP protocol family
 - IP (Internet Protocol) :
 - Provides *basic naming scheme* and unreliable *delivery capability* of packets (datagrams) from *host-to-host*
 - IP mechanism is unreliable in the sense that it makes no effort to recover if datagrams are lost or duplicated in the network.
 - UDP (Unreliable Datagram Protocol)
 - extends IP slightly to provide *unreliable* datagram delivery from *process-to-process*
 - TCP (Transmission Control Protocol)
 - Uses IP to provide *reliable* byte streams from *process-to-process* over *connections*
- Accessed via a mix of Unix file I/O and functions from the *sockets interface*

UDP

FUNCTIONING OF USER DATAGRAM PROTOCOL (UDP)

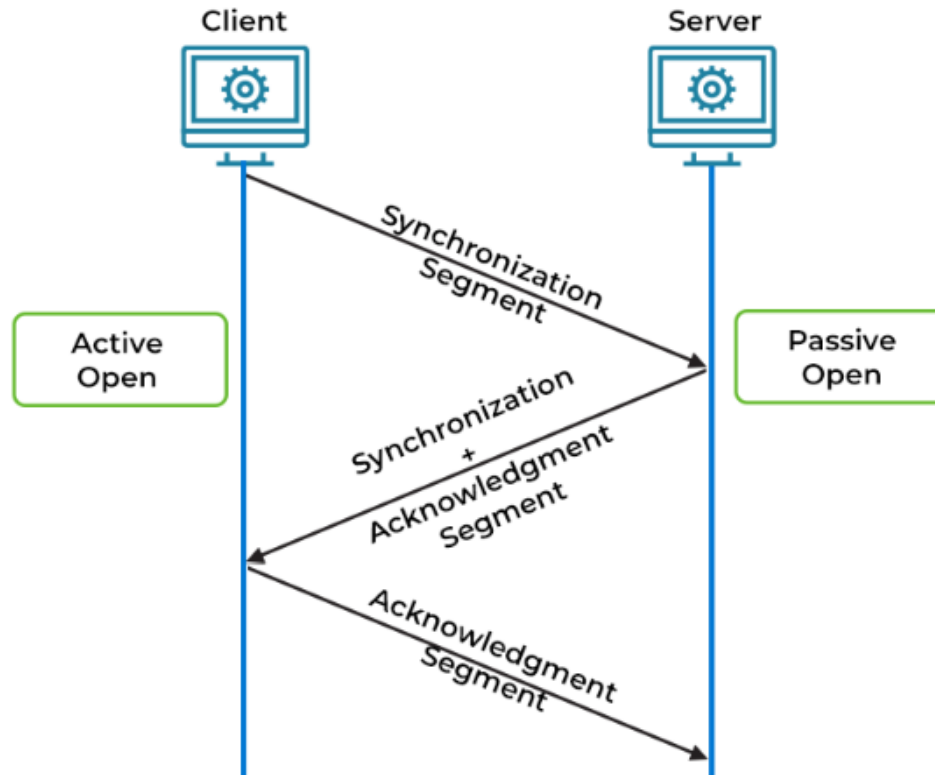


UDP enables continuous data transmission (i.e., response) without acknowledging or confirming the connection

<https://www.spiceworks.com/tech/networking/articles/tcp-vs-udp/>

TCP

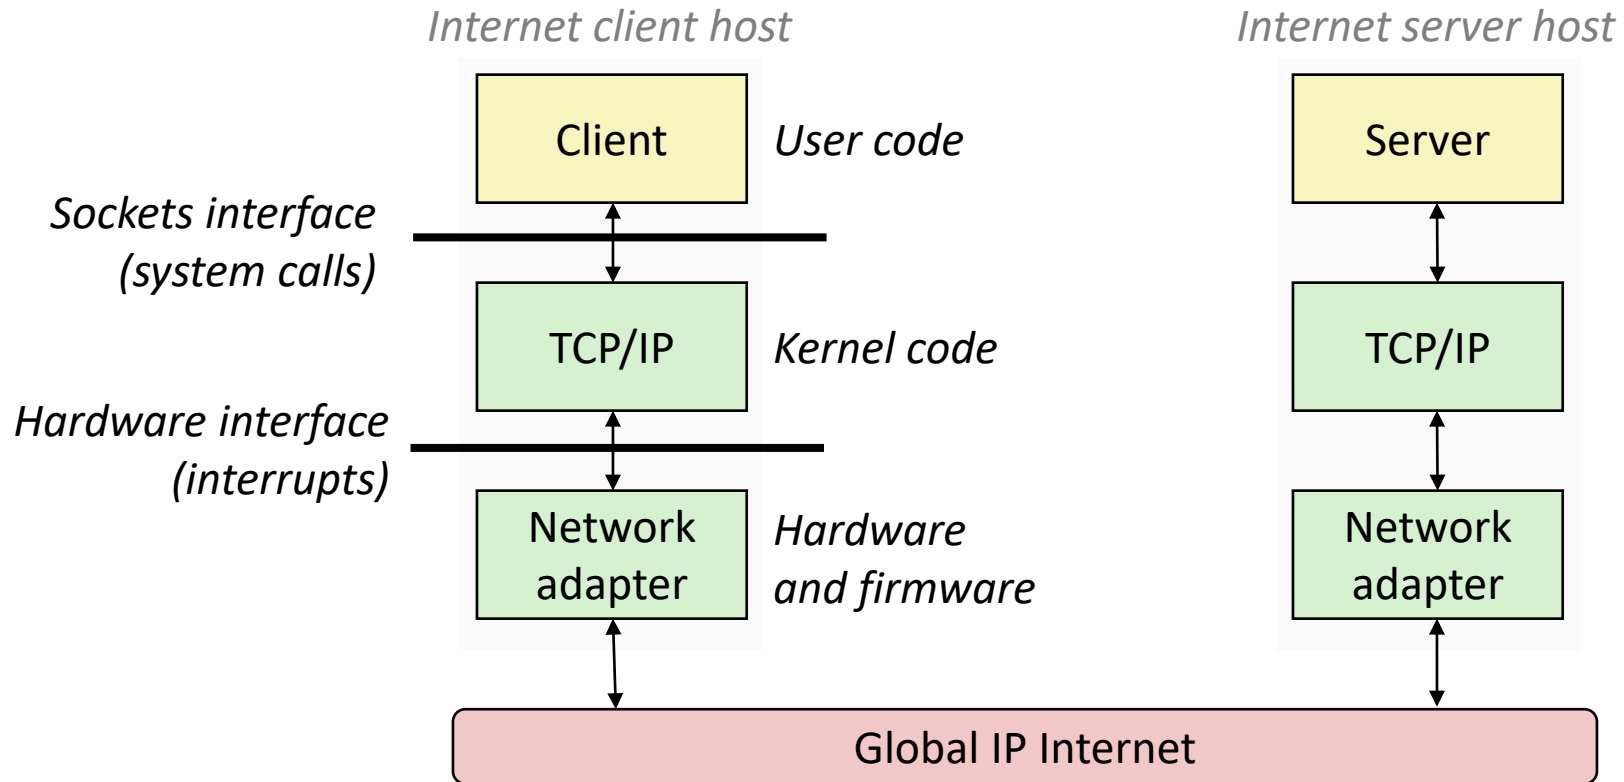
FUNCTIONING OF TRANSMISSION CONTROL PROTOCOL (TCP)



TCP relies on a three-way handshake (synchronization, synchronization acknowledgment, and final acknowledgment)

<https://www.spiceworks.com/tech/networking/articles/tcp-vs-udp/>

Hardware and Software Organization of an Internet Application



- Each Internet host runs software that implements the TCP/IP protocol.
- Internet client and servers communicate using a mix of socket interface functions and Unix I/O functions.
- The socket functions are typically implemented as system call.

A Programmer's View of the Internet

1. Hosts are mapped to a set of 32-bit *IP addresses*

- 161.45.162.100

2. The set of IP addresses is mapped to a set of identifiers called Internet *domain names*

- 161.45.162.100 is mapped to `www.cs.mtsu.edu`

3. A process on one Internet host can communicate with a process on another Internet host over a *connection*

Aside: IPv4 and IPv6

- The original Internet Protocol, with its 32-bit addresses, is known as *Internet Protocol Version 4* (**IPv4**)
- 1996: Internet Engineering Task Force (IETF) introduced *Internet Protocol Version 6* (**IPv6**) with 128-bit addresses
 - Intended as the successor to IPv4
- As of 2022, vast majority of Internet traffic still carried by IPv4
 - Only around 38% of users access Google services using IPv6.
- We will focus on IPv4, but will show you how to write networking code that is protocol-independent.

(1) IP Addresses

- An IP address is an unsigned 32-bit integer.
- Network program stores IP addresses in the *IP address structure (shown below)*
 - IP addresses are always stored in memory in *network byte order* (big-endian byte order)
 - Because Internet hosts can have different host byte order, TCP/IP defines a uniform network byte order (big-endian byte order) for any integer data item, such as IP address, that is carried across the network in a packet header.
 - Next slides shows a list of Unix functions for converting network and host byte order.

```
/* Internet address structure */  
struct in_addr {  
    uint32_t    s_addr; /* network byte order (big-endian) */  
};
```


Unix Functions for byte order conversion

Useful network byte-order conversion functions

("l" = 32 bits and "s" = 16 bits)

(n => network, h => host)

htonl: convert uint32_t from host to network byte order

htons: convert uint16_t from host to network byte order

ntohl: convert uint32_t from network to host byte order

ntohs: convert uint16_t from network to host byte order

```
#include <arpa/inet.h>
//Return value in network byte order
uint32_t htonl(uint32_t host);
uint16_t htons(uint16_t host);

//Return value in host byte order
uint32_t ntohl(uint32_t net);
uint16_t ntohs(uint16_t net);
```

Dotted Decimal Notation

- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
 - IP address: `0x8002C2F2` = `128.2.194.242`
- Use `getaddrinfo` and `getnameinfo` functions (described later) to convert between IP addresses and dotted decimal format.

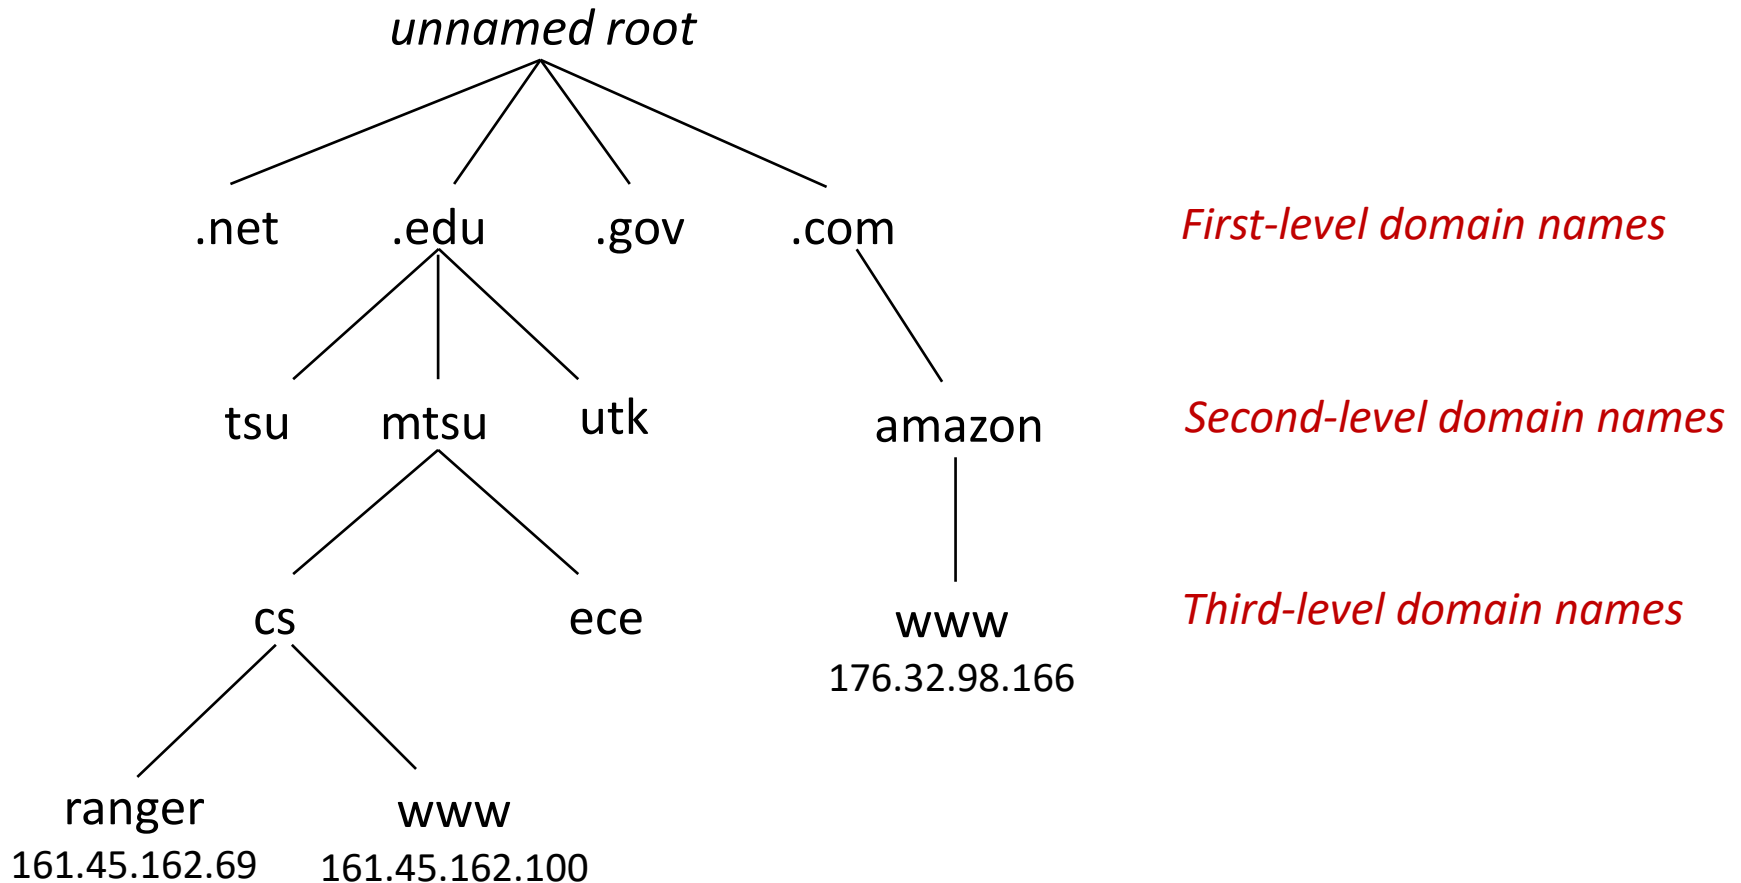
Exercise

- **Convert the following 32-bit addresses into dotted decimal notation.**

- 0xffffffff

- 0x0aff090c

(2) Internet Domain Names



History

- **Before Domain Name System(DNS), all mappings were in `hosts.txt`**
 - `/etc/hosts` on linux
 - `C:\Windows\System32\drivers\etc\hosts` on windows
- **Centralized, manual system**
 - SRI was the first organization to assign website address such as www.sri.com with extensions such as “.com”, “.org”.
 - These addresses were assigned to network hosts by the Network Information Center (NIC), managed by SRI from 1970 until 1991.
 - Changes were submitted to SRI via email
 - Machines periodically FTP new copies of *hosts.txt*
 - Administrators could pick names at their discretion
 - Any name was allowed

Towards DNS

- **Eventually, the *hosts.txt* system fell apart**
 - Not scalable, SRI couldn't handle the load
 - Hard to enforce uniqueness of names
 - e.g MIT
 - Massachusetts Institute of Technology?
 - Melbourne Institute of Technology?
 - Many machines had inaccurate copies of *hosts.txt*
- **Thus, DNS was born**

Domain Naming System (DNS)

- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called *DNS*
 - No Centralization
- Conceptually, programmers can view the DNS database as a collection of millions of *host entries*.
 - Each host entry defines the mapping between a set of domain names and IP addresses.
- Hierarchical namespace
 - As opposed to original, flat namespace
 - E.g. .com -> google.com -> mail.google.com

Properties of DNS Mappings

- Can explore properties of DNS mappings using `nslookup`
 - Output edited for brevity
- Each host has a locally defined domain name `localhost` which always maps to the *loopback address* `127.0.0.1`

```
linux> nslookup localhost  
Address: 127.0.0.1
```

- Use `hostname` to determine real domain name of local host:

```
linux> hostname  
csci3240-00.cs.mtsu.edu
```


Properties of DNS Mappings (cont)

- **Simple case: one-to-one mapping between domain name and IP address:**

```
linux> nslookup csci3240-00.cs.mtsu.edu  
Address: 161.45.164.116
```

- **Multiple domain names mapped to the same IP address:**

```
linux> nslookup www.mtsu.edu  
Address: 10.14.0.116  
linux> nslookup w1.mtsu.edu  
Address: 10.14.0.116
```

Properties of DNS Mappings (cont)

- **Multiple domain names mapped to multiple IP addresses:**

```
linux> nslookup www.twitter.com
Address: 199.16.156.6
Address: 199.16.156.70
Address: 199.16.156.102
Address: 199.16.156.230

linux> nslookup twitter.com
Address: 199.16.156.102
Address: 199.16.156.230
Address: 199.16.156.6
Address: 199.16.156.70
```

- **Some valid domain names don't map to any IP address:**

```
linux> nslookup csc.mtsu.edu
*** Can't find csc.mtsu.edu : No answer
```

(3) Internet Connections

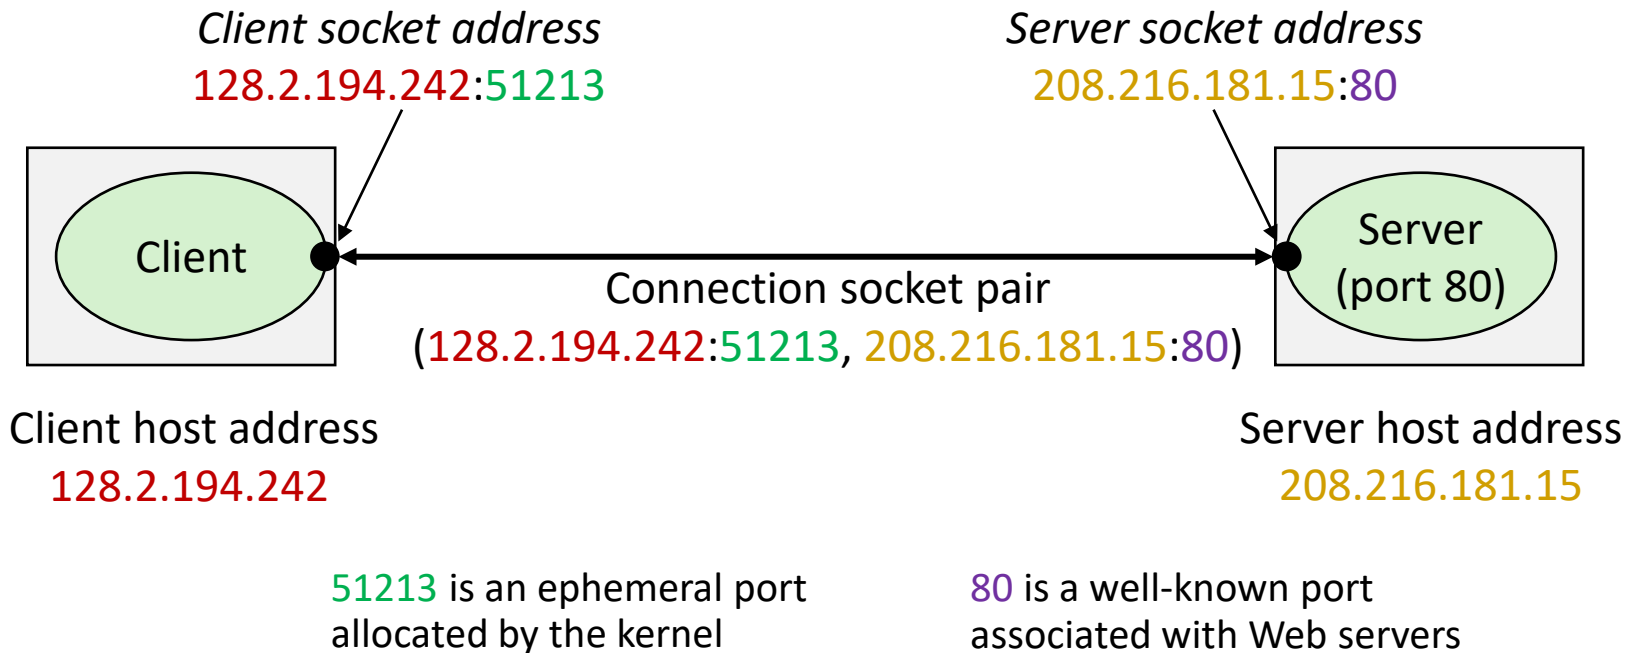
- Clients and servers communicate by sending streams of bytes over **connections**. Each connection is:
 - *Point-to-point*: connects a pair of processes.
 - *Full-duplex*: data can flow in both directions at the same time,
 - *Reliable*: stream of bytes sent by the source is eventually received by the destination in the same order it was sent.
- A **socket** is an endpoint of a connection
 - *Socket address* is an **IPAddress:port** pair
- A **port** is a 16-bit integer that identifies a process:
 - **Ephemeral port**: Assigned automatically by client kernel when client makes a connection request.
 - **Well-known port**: Associated with some **service** provided by a server (e.g., port 80 is associated with Web servers)

Well-known Ports and Service Names

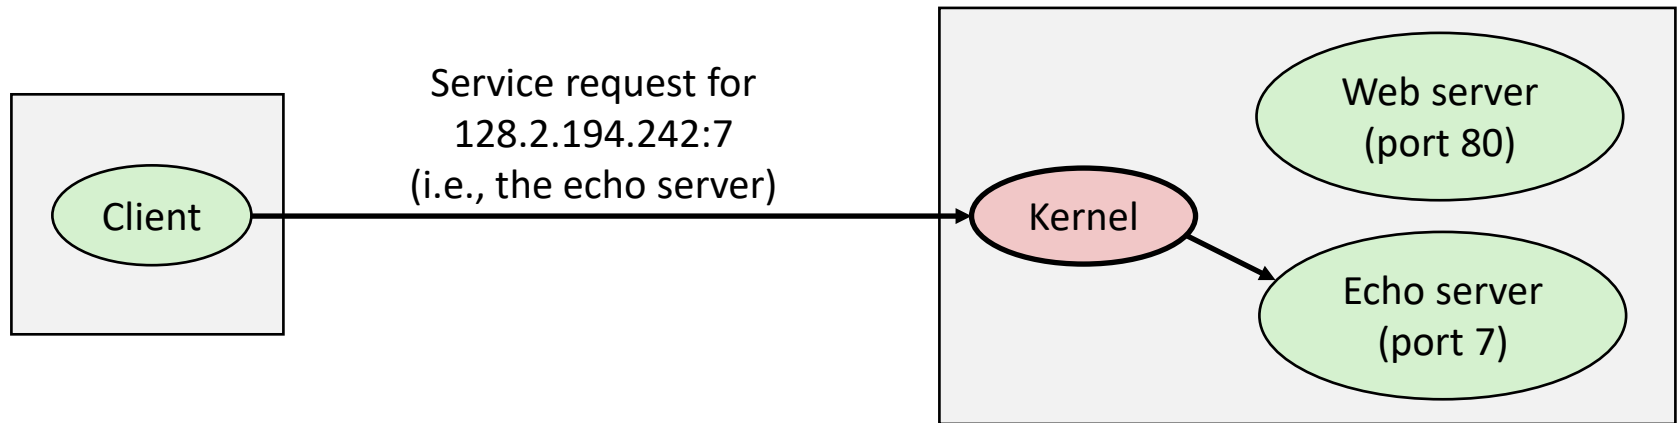
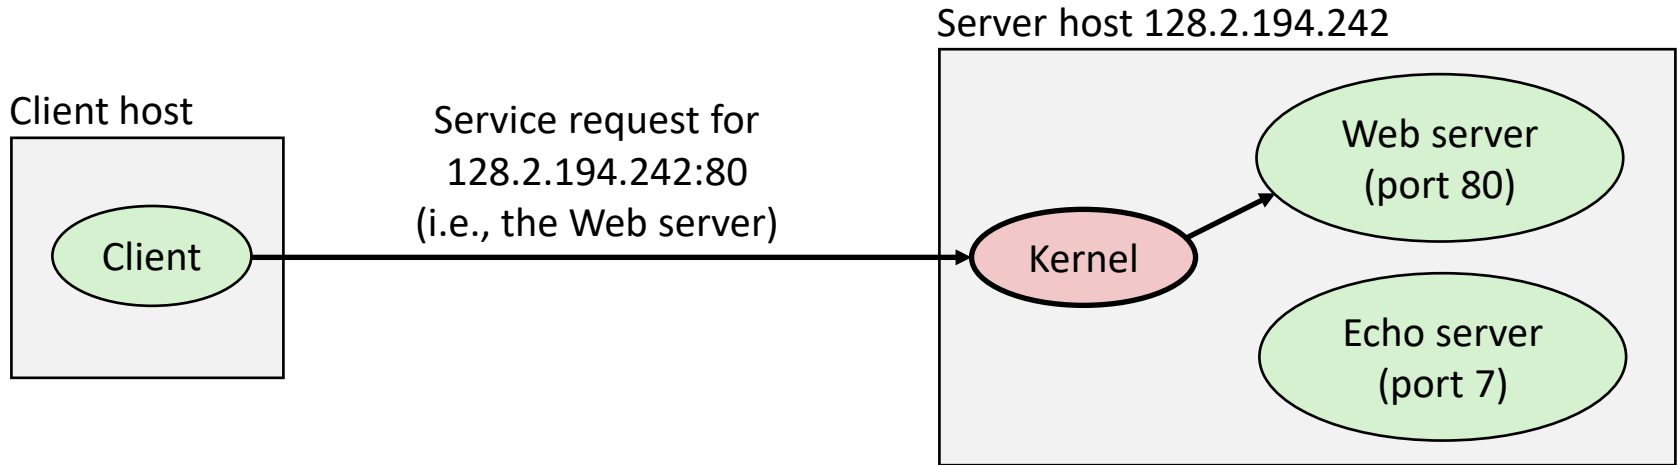
- Popular services have permanently assigned *well-known ports* and corresponding *well-known service names*:
 - echo server: 7/echo
 - ssh servers: 22/ssh
 - email server: 25/smtp
 - Web servers: 80/http
- Mappings between well-known ports and service names is contained in the file `/etc/services` on each Linux machine.

Anatomy of a Connection

- A connection is uniquely identified by the socket addresses of its endpoints (*socket pair*)
 - (cliaddr:cliport, servaddr:servport)



Using Ports to Identify Services



Sockets Interface

- **Set of system-level functions are used in conjunction with Unix I/O to build network applications.**
- **Available on all modern systems**
 - Unix variants, Windows, OS X, IOS, Android, ARM

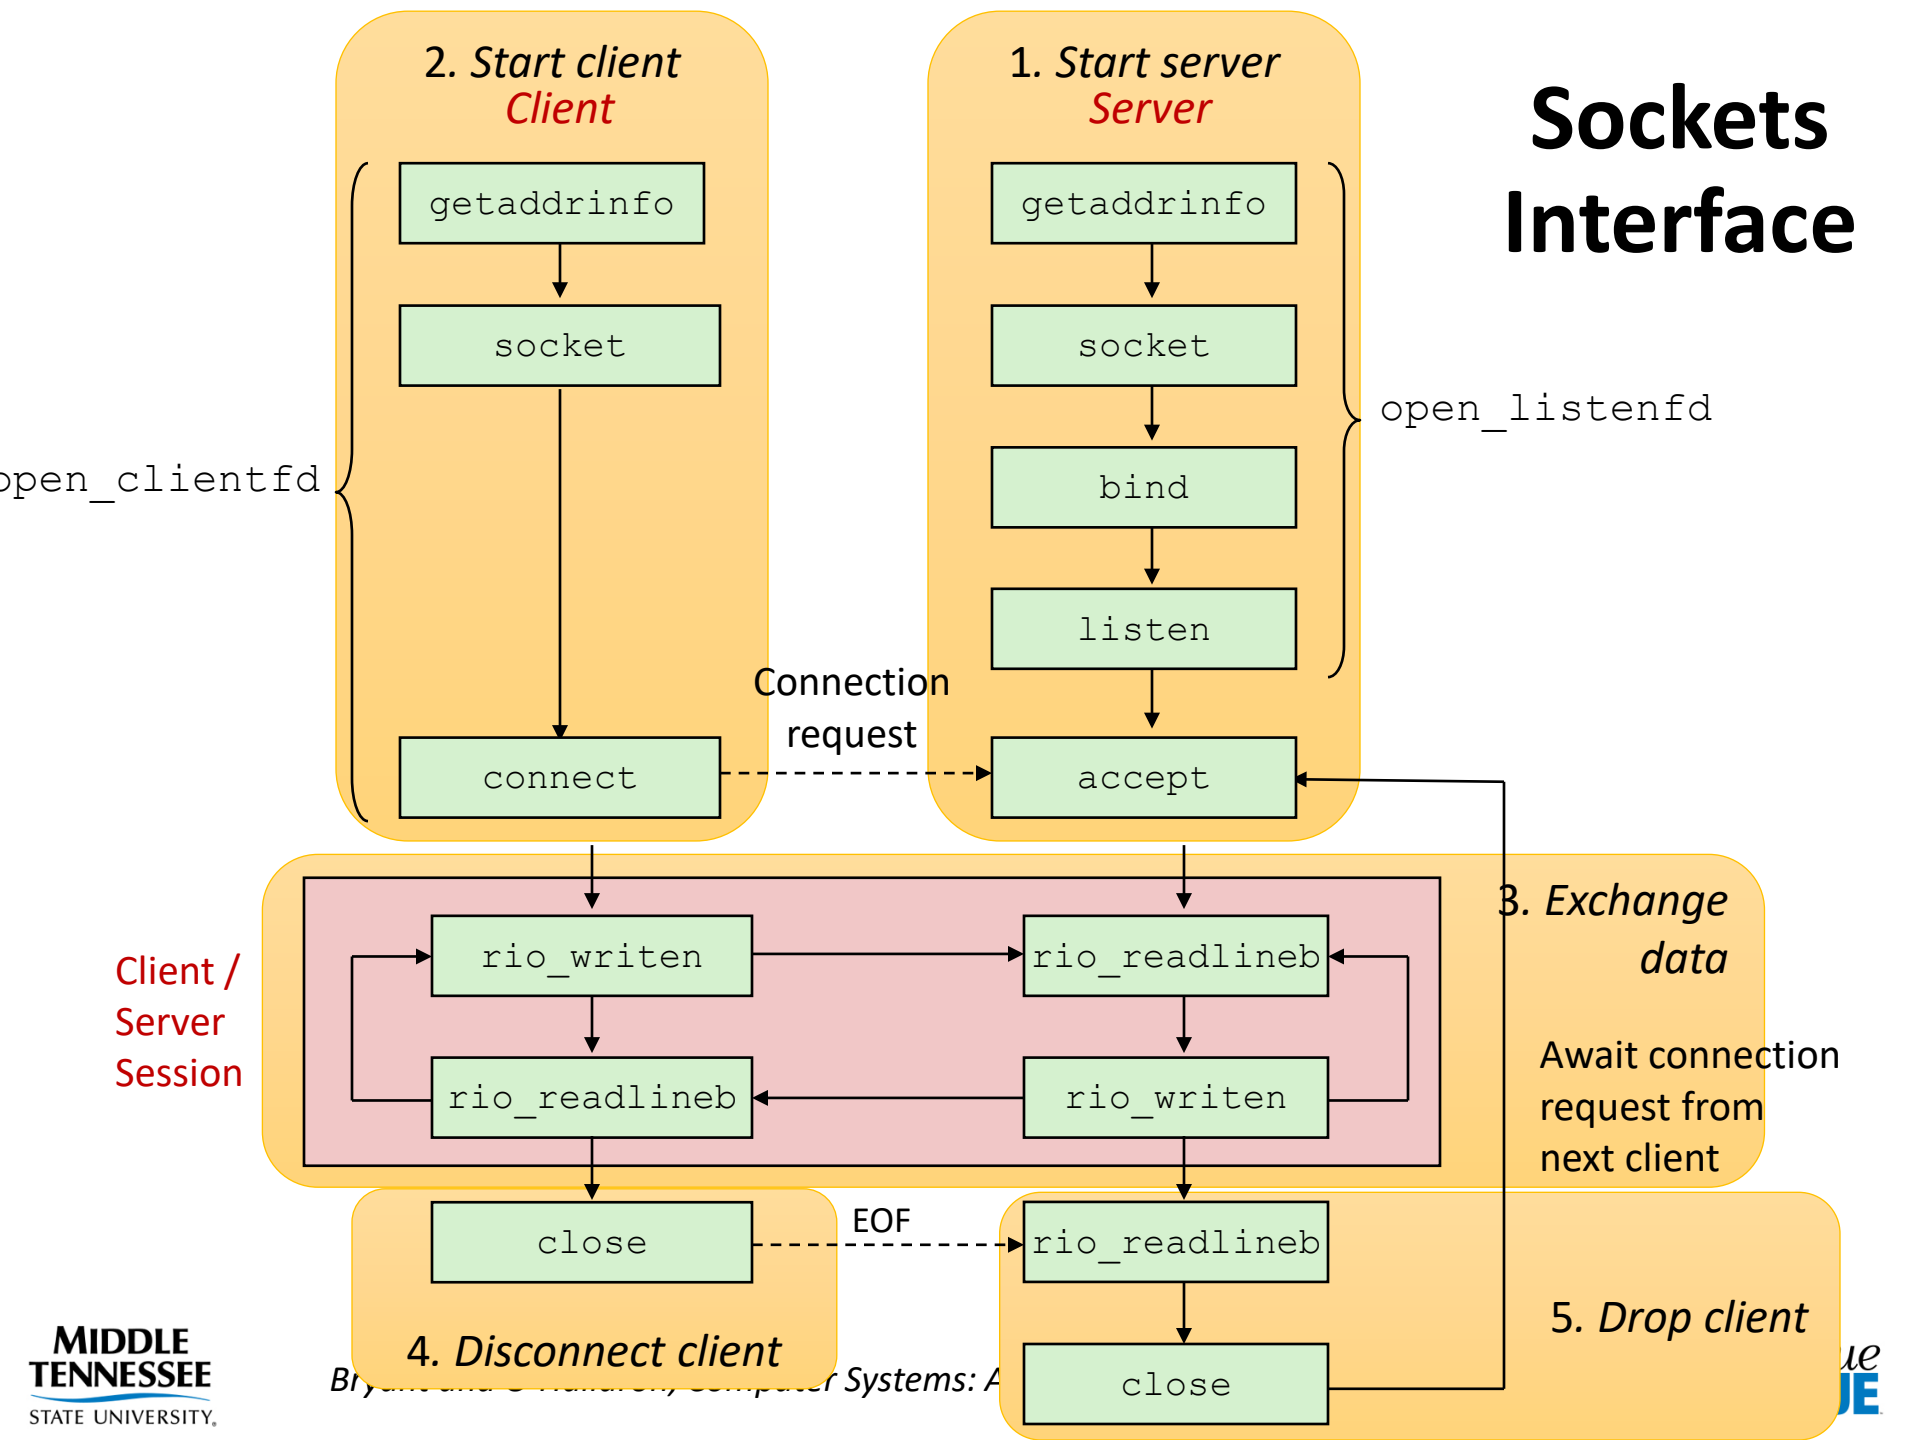
Sockets

- **What is a socket?**
 - To the kernel, a socket is an endpoint of communication
 - To an application (program), a socket is a file descriptor that lets the application read/write from/to the network
 - **Remember:** All Unix I/O devices, including networks, are modeled as files
- **Clients and servers communicate with each other by reading from and writing to socket descriptors**



- **The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors**

Sockets Interface



Sockets Address Structure

- Internet socket addresses are stored in 16-byte structures having the type **sockaddr_in**
- For Internet applications, **sin_family** field is AF_INET
- The **sin_port** field is a 16-bit port number
- And, the **sin_addr** field contains a 32-bit IP address.
- The IP address and port number are always stored in network byte order

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```

*Here, the **_in** suffix in **sockaddr_in** refers to internet, not input.*

Socket Address Structures

▪ Generic socket address:

- The **connect**, **bind**, and **accept** functions requires a pointer to a protocol specific socket address structure.
- The problem faced by the designer of the sockets interface was how to define these functions to accept any kind of socket address structure.
- Necessary only because C did not have generic (**void ***) pointers when the sockets interface was designed.
- The solution was to define sockets function to expect a pointer to a generic sockaddr structure and then require applications to cast any pointers to protocol-specific structures to this generic structure.
- For casting convenience, we can use typedef:

typedef struct sockaddr SA;

```
struct sockaddr {  
    uint16_t    sa_family;    /* Protocol family */  
    char        sa_data[14];  /* Address data. */  
};
```

sa_family



Family Specific

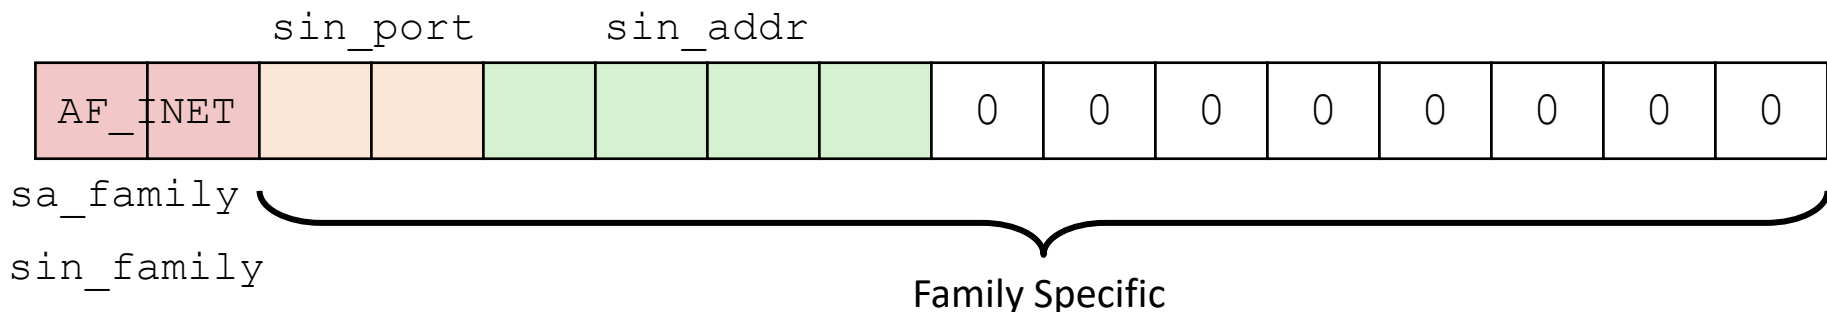
Socket Address Structures

■ Internet-specific socket address:

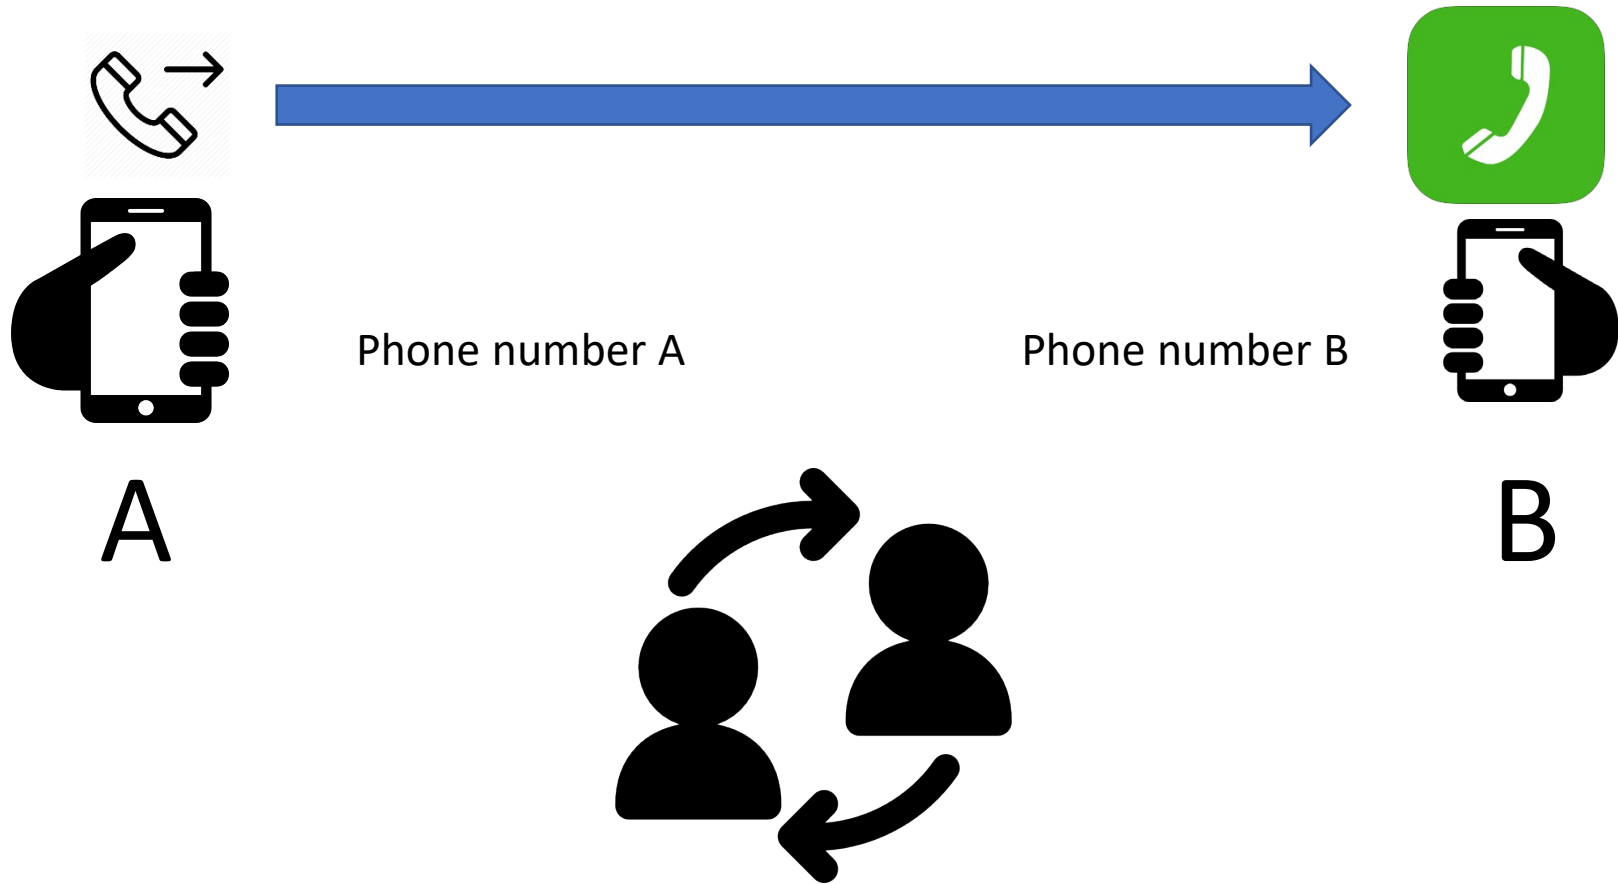
- Must cast `(struct sockaddr_in *)` to `(struct sockaddr *)` for functions that take socket address arguments.

```
struct sockaddr_in {
    uint16_t      sin_family; /* Protocol family (always AF_INET) */
    uint16_t      sin_port;   /* Port num in network byte order */
    struct in_addr sin_addr;   /* IP addr in network byte order */
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */
};
```

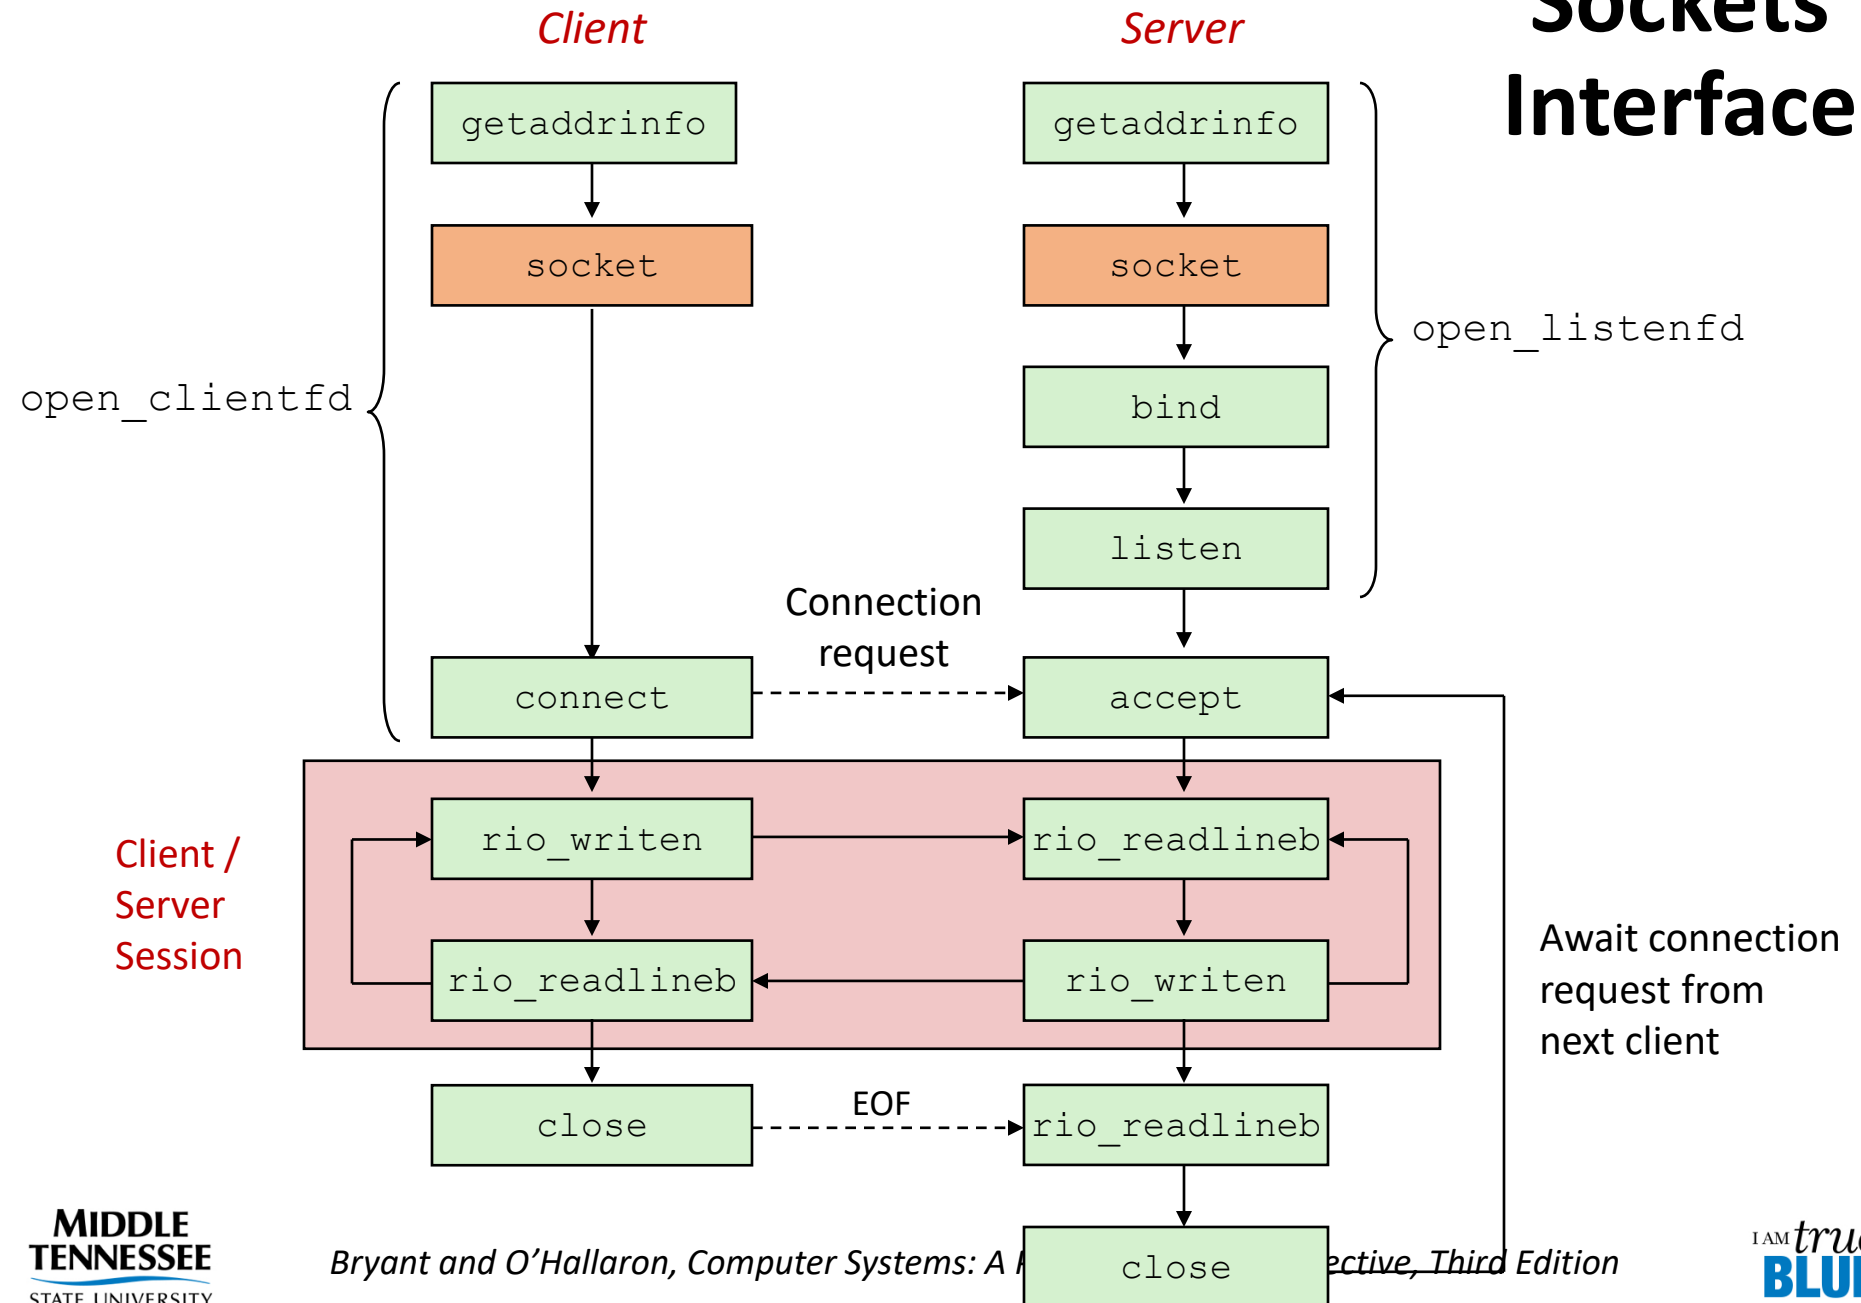
```
struct sockaddr {
    uint16_t  sa_family; /* Protocol family */
    char      sa_data[14]; /* Address data. */
};
```



Socket Programming: An Analogy with a Phone call



Sockets Interface



Sockets Interface: `socket`

- TCP Socket can be compared to a phone
 - We need a device to call someone or accept their phone calls, right?
- To perform network I/O, the first thing a process must do is, call the `socket` function to create a specific type of socket by specifying the type of communication protocol desired, protocol family, etc.

Sockets Interface: `socket`

- Clients and servers use the `socket` function to create a *socket descriptor*.

```
#include <sys/socket.h>
int socket(int domain, int type, int protocol)
```

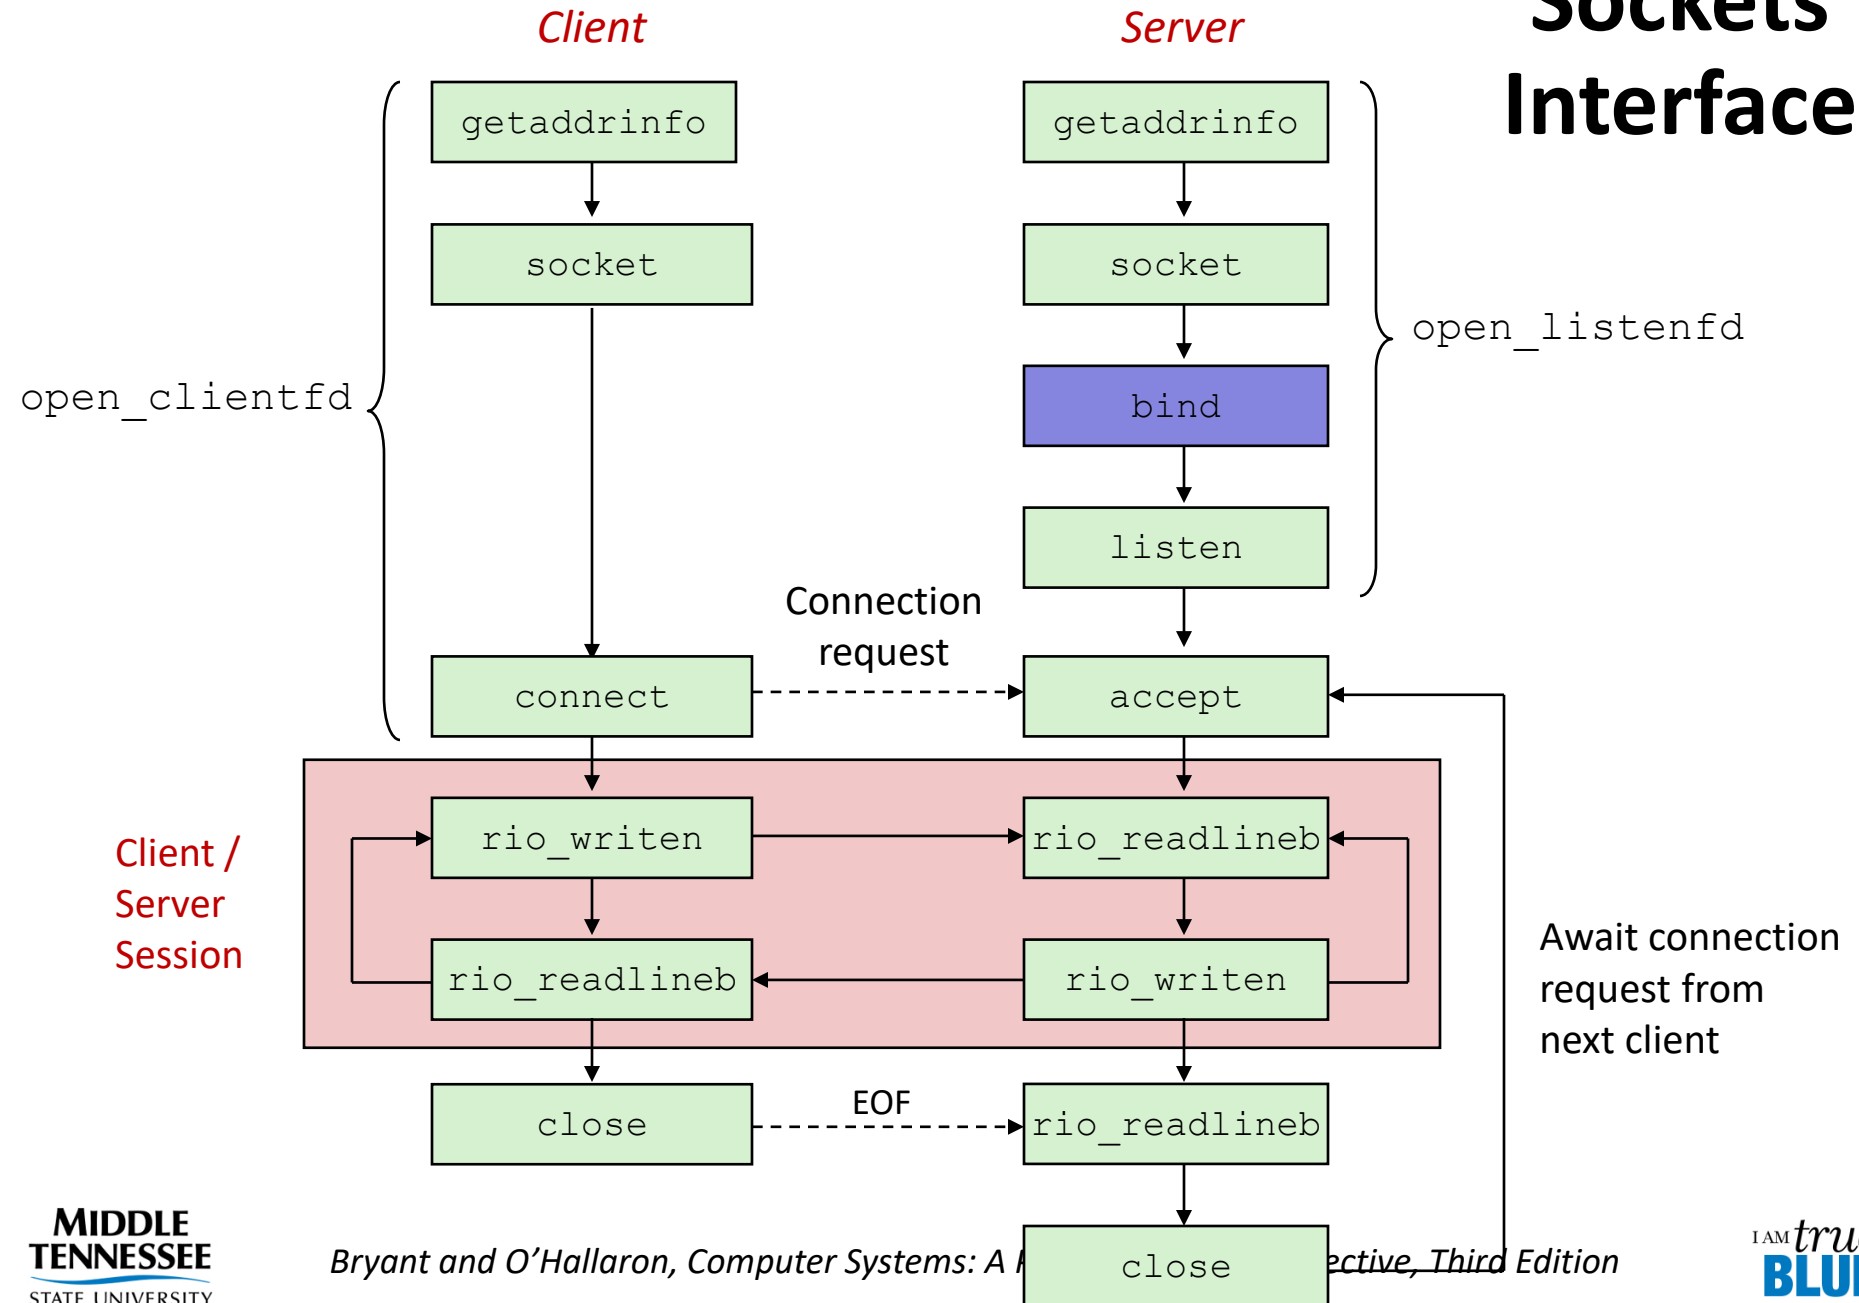
- `int clientfd = Socket(AF_INET, SOCK_STREAM, 0);`

Indicates that we are using
32-bit IPV4 addresses

Indicates that the socket
will be the end point of a
connection

*Protocol specific! Best practice is to use **getaddrinfo** to generate the parameters automatically, so that code is protocol independent.*

Sockets Interface



Sockets Interface: bind

- It is like registering a phone number to the phone (to connect to other phones).
- Sockets do not have a complete address in the beginning to start transferring the data, so we bind a Socket to a port.
 - Analogy: Phones do not have a phone number at the beginning to make phone calls, so we register a phone number.
- The process of allocating a port number to a socket is called 'binding'.

Sockets Interface: `bind`

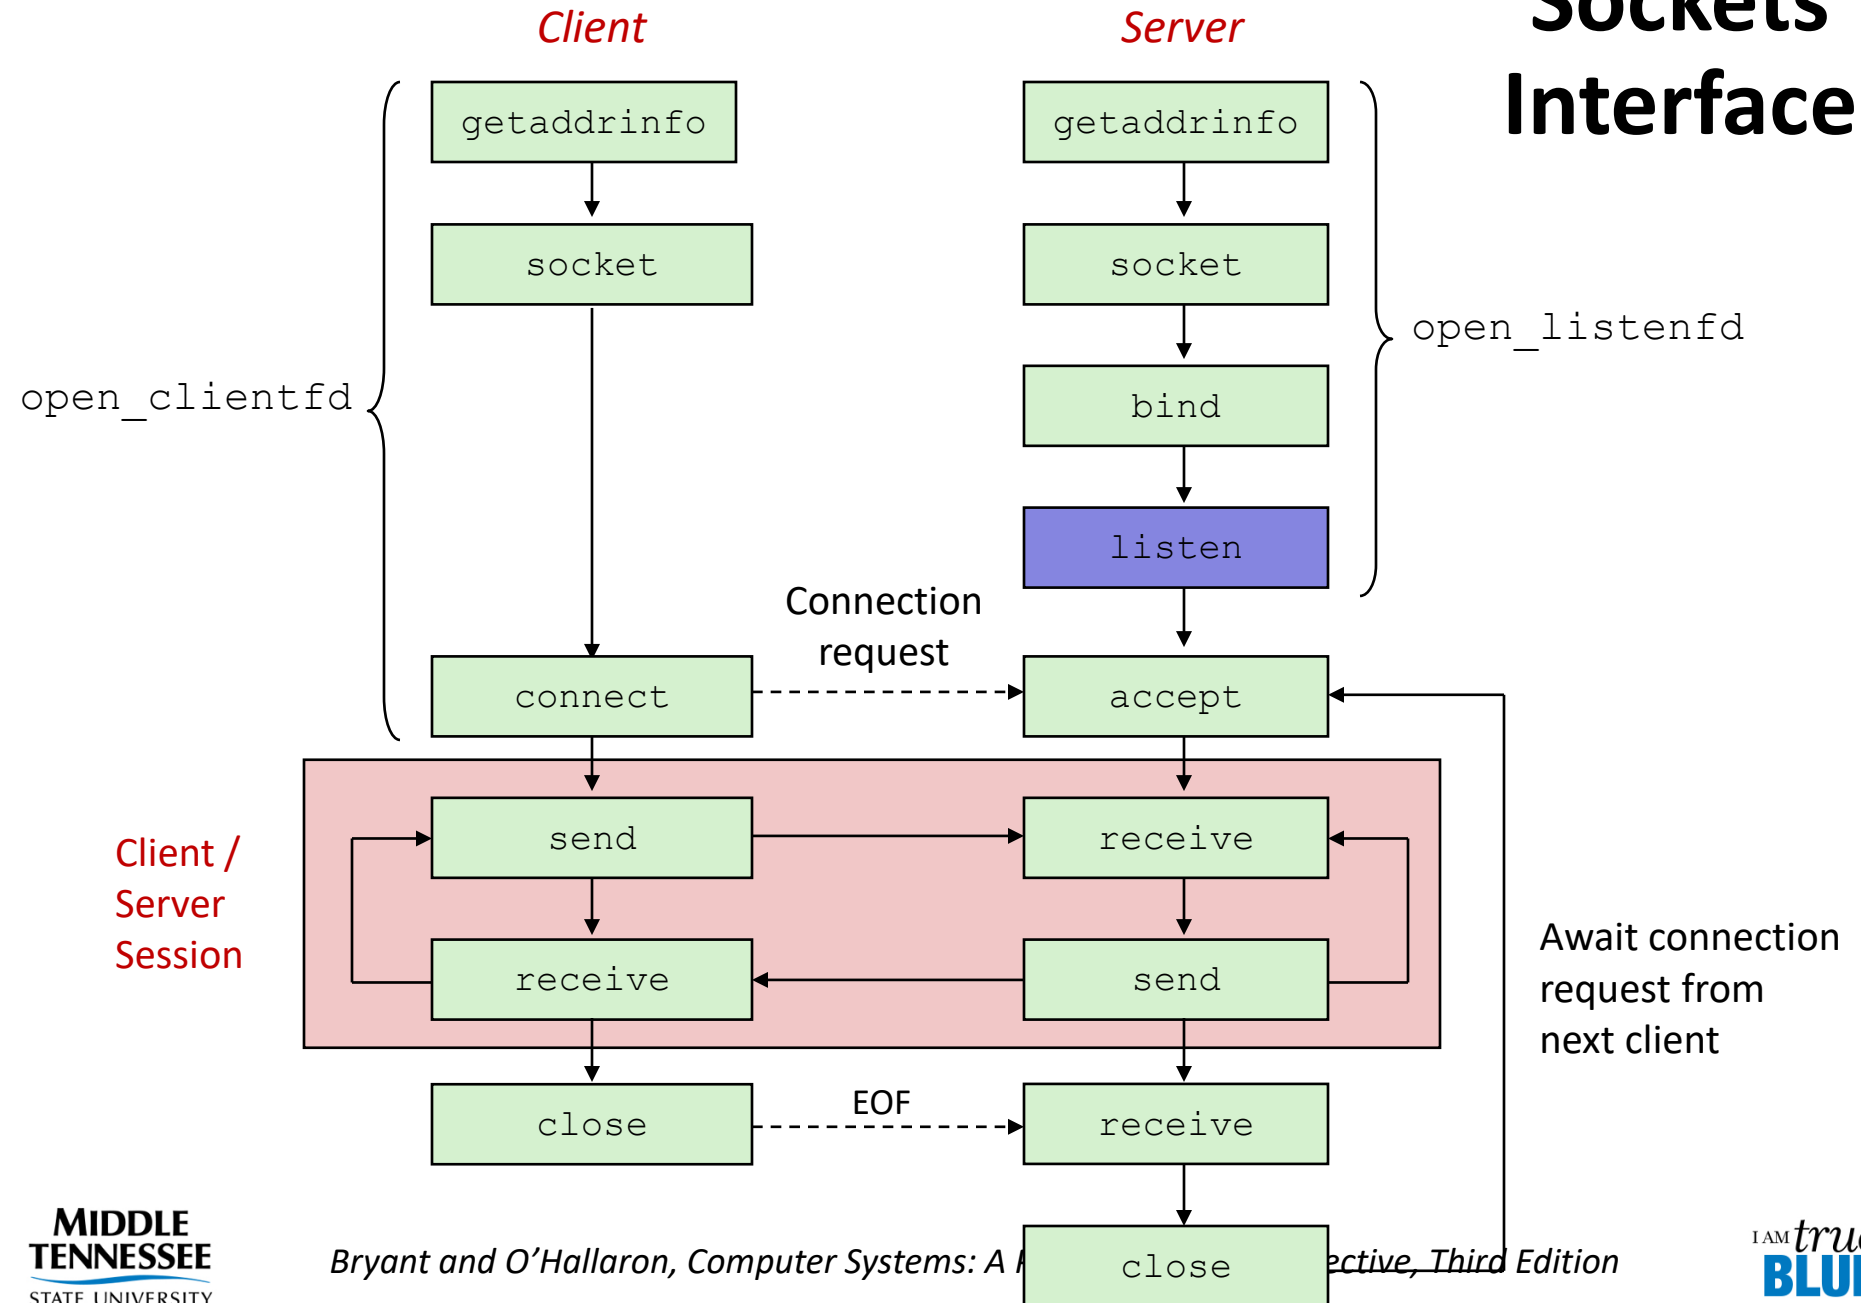
- A server uses `bind` to ask the kernel to associate the server's socket address with a **socket descriptor**:

```
#include <sys/socket.h>
int bind(int sockfd, SA *addr, socklen_t addrlen);
```

- The process can read bytes that arrive on the connection whose endpoint is **addr** by reading from descriptor **sockfd**.
- Similarly, writes to **sockfd** are transferred along connection whose endpoint is **addr**.

*Best practice is to use **getaddrinfo** to supply the arguments **addr** and **addrlen**.*

Sockets Interface



Sockets Interface: `listen`

- Clients are active entities that initiate connection requests.
- Servers are passive entities that wait for connection requests from clients.
- By default, the kernel assumes that a descriptor created by the socket function corresponds to an *active* socket that will be live on the client end of a connection.
- A server calls the listen function to tell the kernel that a descriptor will be used by a server rather than a client:

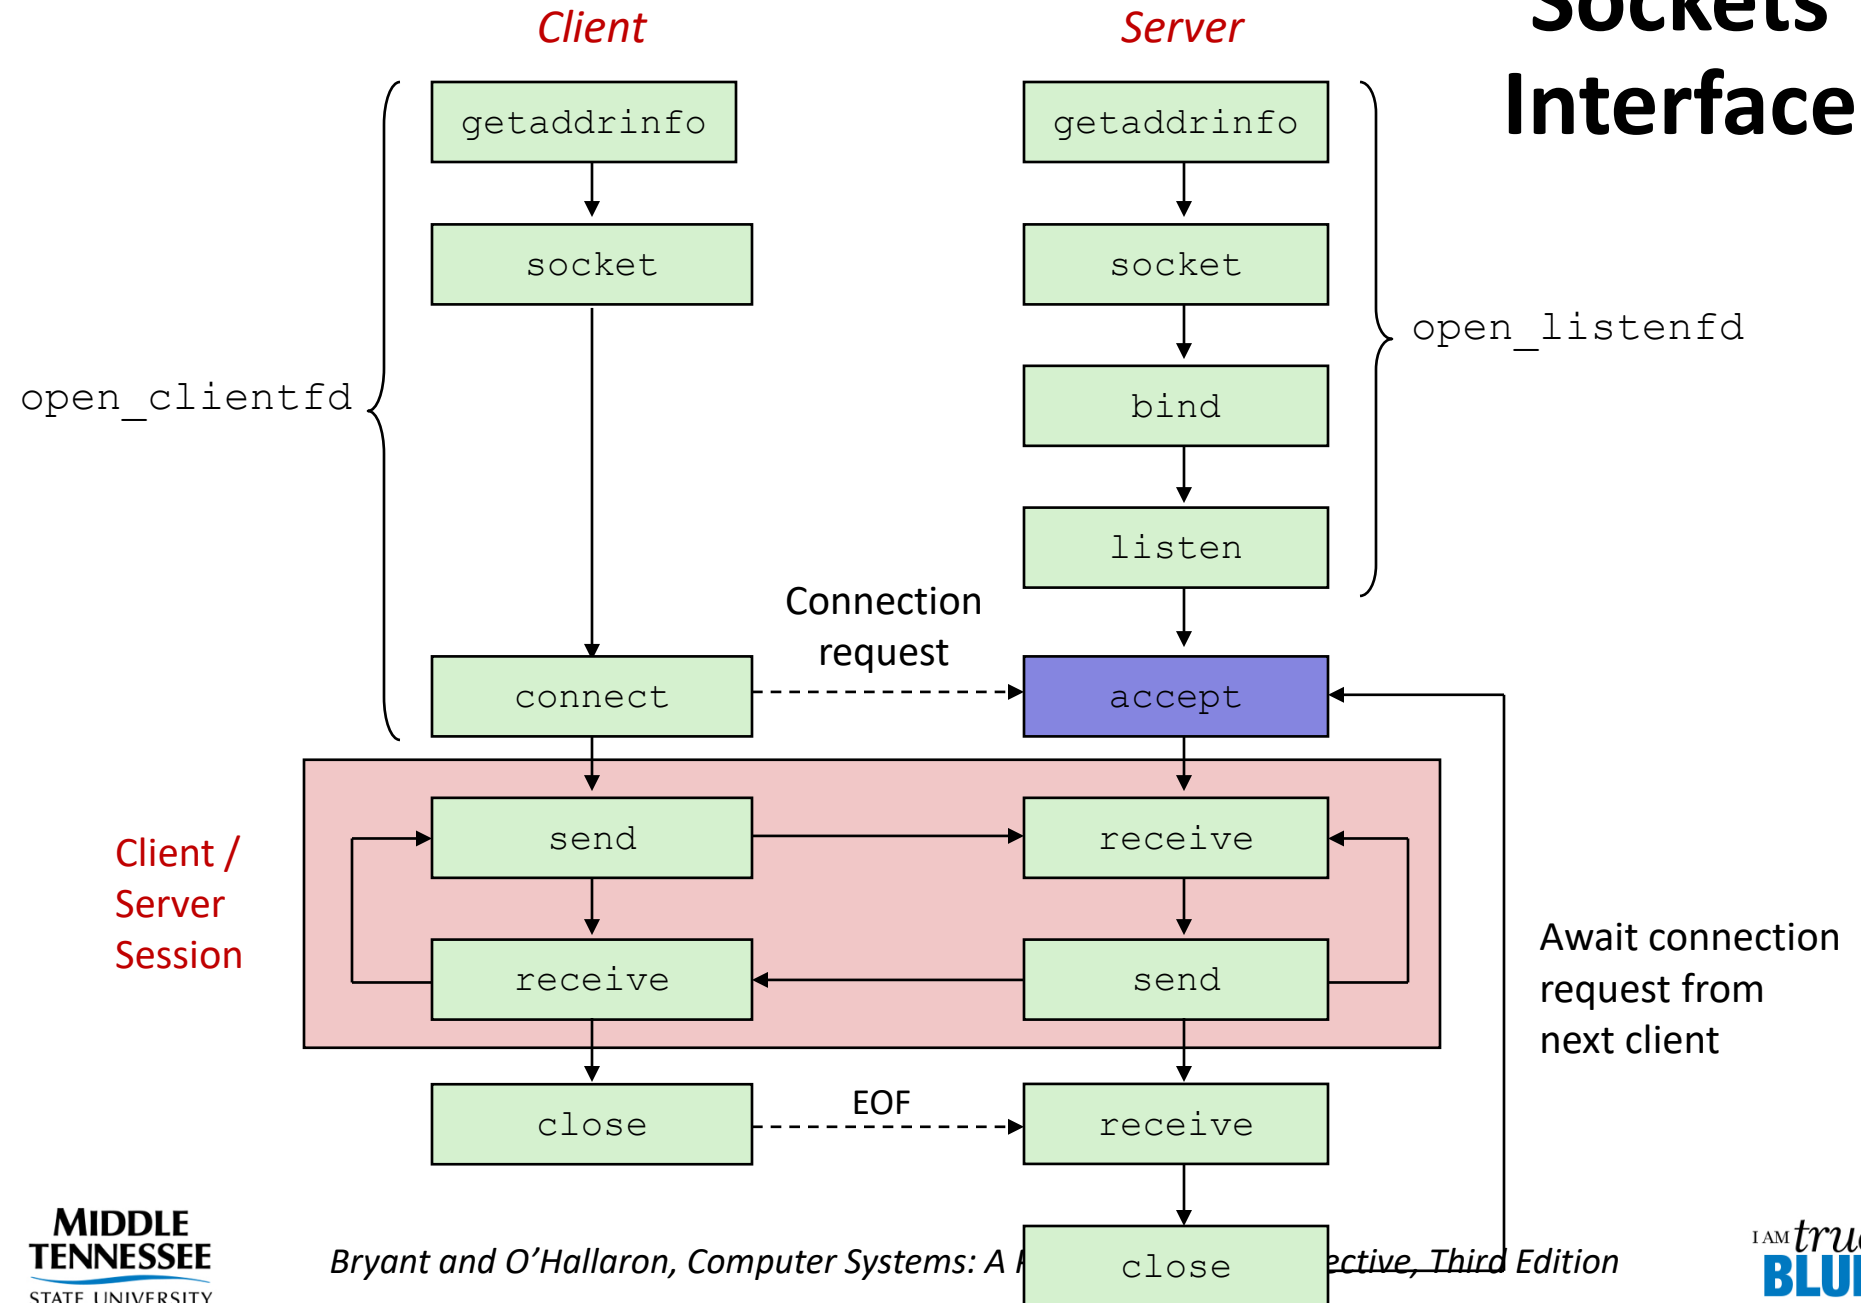
```
int listen(int sockfd, int backlog);
```

Sockets Interface: `listen`

```
int listen(int sockfd, int backlog);
```

- The `listen` function converts `sockfd` from an active socket to a *listening socket* that can accept connection requests from clients.
- The backlog is a hint about the number of outstanding connection requests that the kernel should queue up before starting to refuse requests.
- We will typically set it to a large value, such as 1024.

Sockets Interface



Sockets Interface: `accept`

- Accept a connection request
 - Transition of the connection request from `listen()` method to an actual Socket
- Servers wait for connection requests from clients by calling `accept` function:

```
int accept(int listenfd, SA *addr, int *addrlen);
```

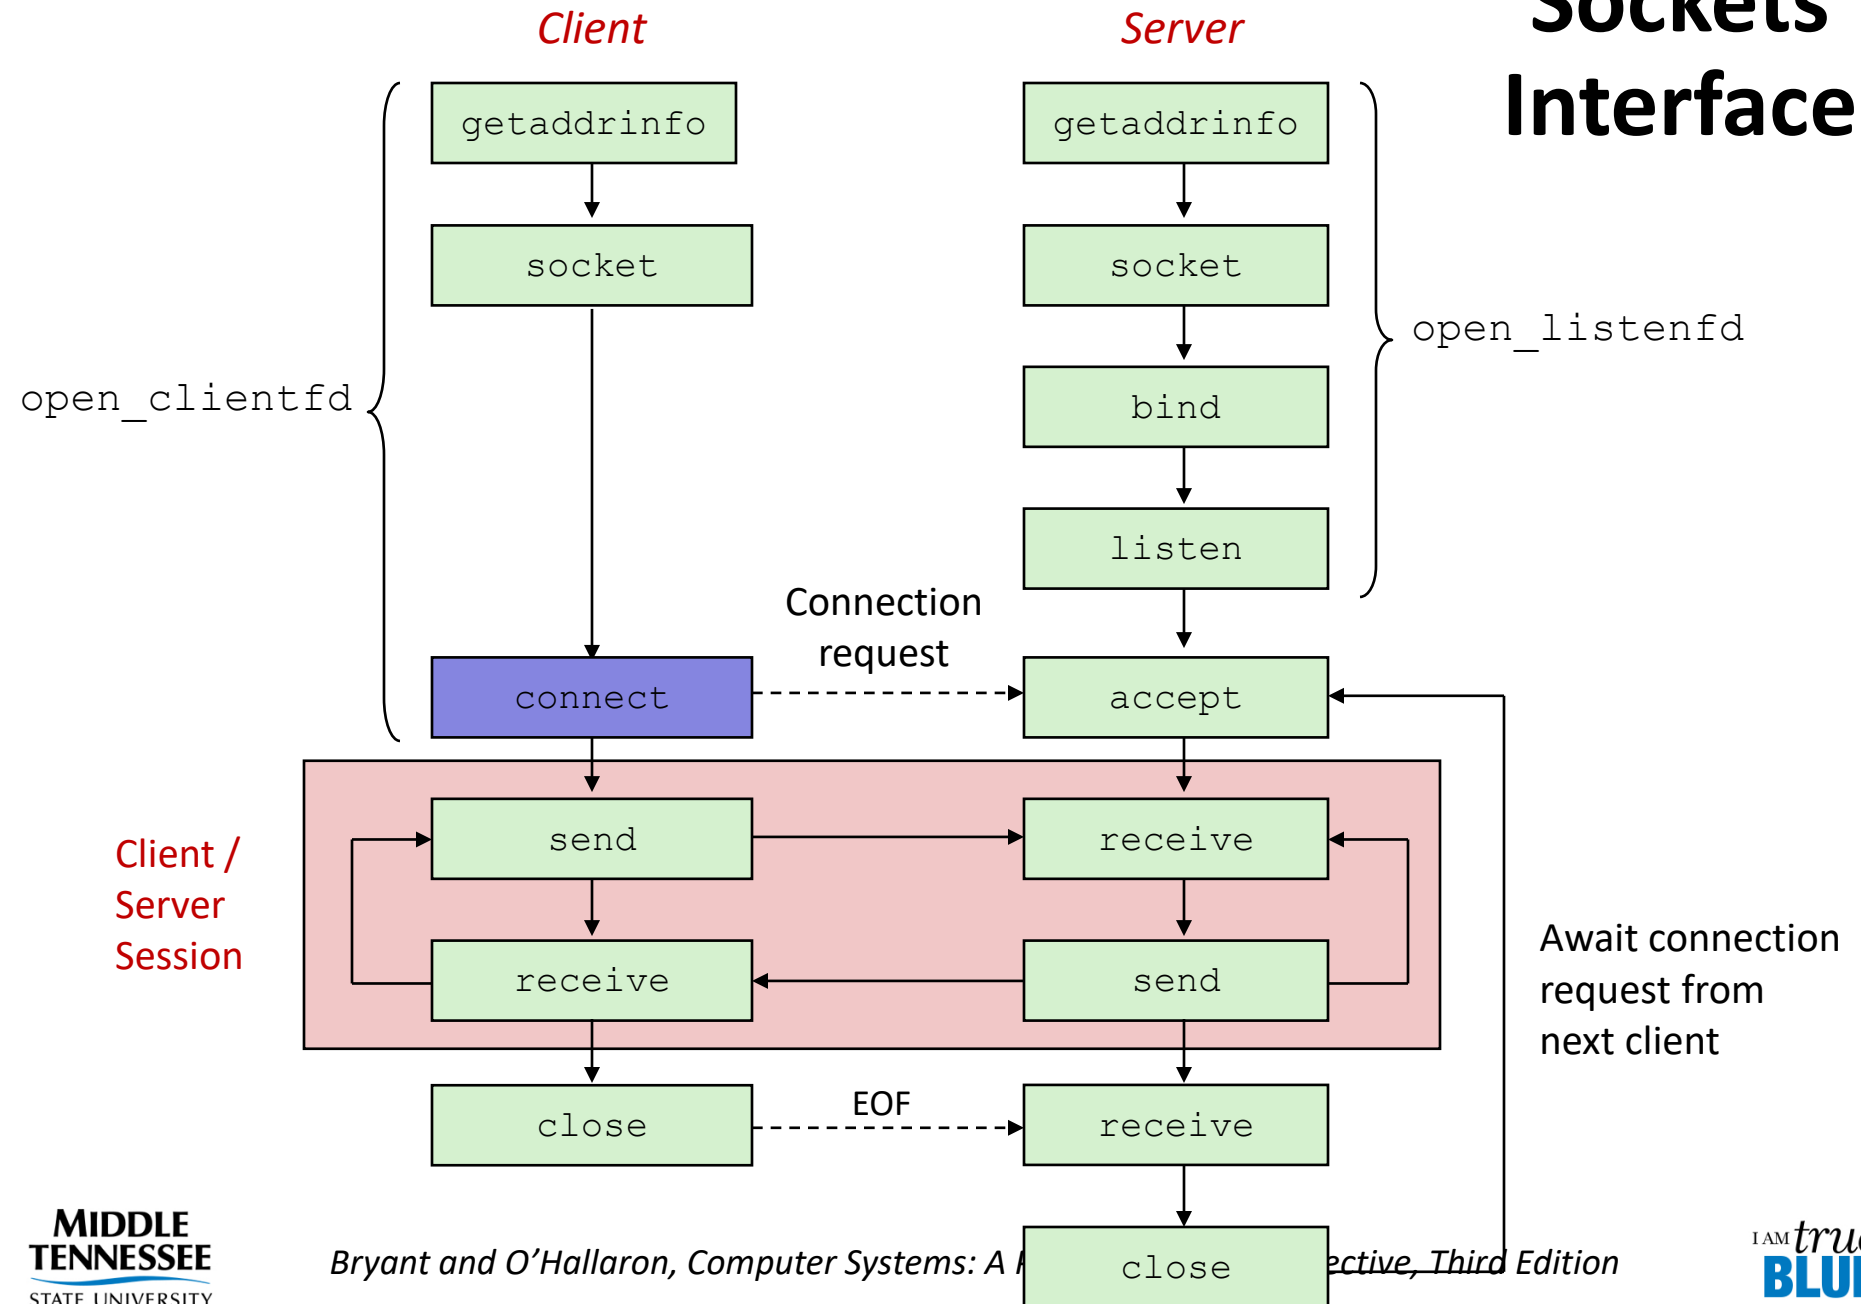
- The `accept` function waits for connection request to arrive on the connection bound to listening descriptor `listenfd`, then fills in client's socket address in `addr` and size of the socket address in `addrlen`.
- Finally, it returns a *connected descriptor* that can be used to communicate with the client via Unix I/O routines.

Make the logical steps from the 4 previous parts

▪ Sequence

- Step 1. Creating the socket (socket function call)
- Step 2. Bind the socket to a port (bind function call)
- Step 3. Listen for incoming connection requests & identify ones (listen function call)
- Step 4. Accept the identified connection request and open the socket (accept function call)

Sockets Interface



Sockets Interface: connect

- A client establishes a connection with a server by calling connect:

```
int connect(int clientfd, SA *addr, socklen_t addrlen);
```

- Attempts to establish a connection with server at socket address `addr`
 - If successful, then `clientfd` is now ready for reading and writing.
 - Resulting connection is characterized by socket pair

`(x:y, addr.sin_addr:addr.sin_port)`


- `x` is client address
- `y` is ephemeral port that uniquely identifies client process on client host

*Best practice is to use **getaddrinfo** to supply the arguments `addr` and `addrlen`.*


accept Illustrated



1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`



2. Client makes connection request by calling and blocking in `connect`



3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`

Connected vs. Listening Descriptors

■ Listening descriptor

- End point for client connection requests
- Created once and exists for lifetime of the server

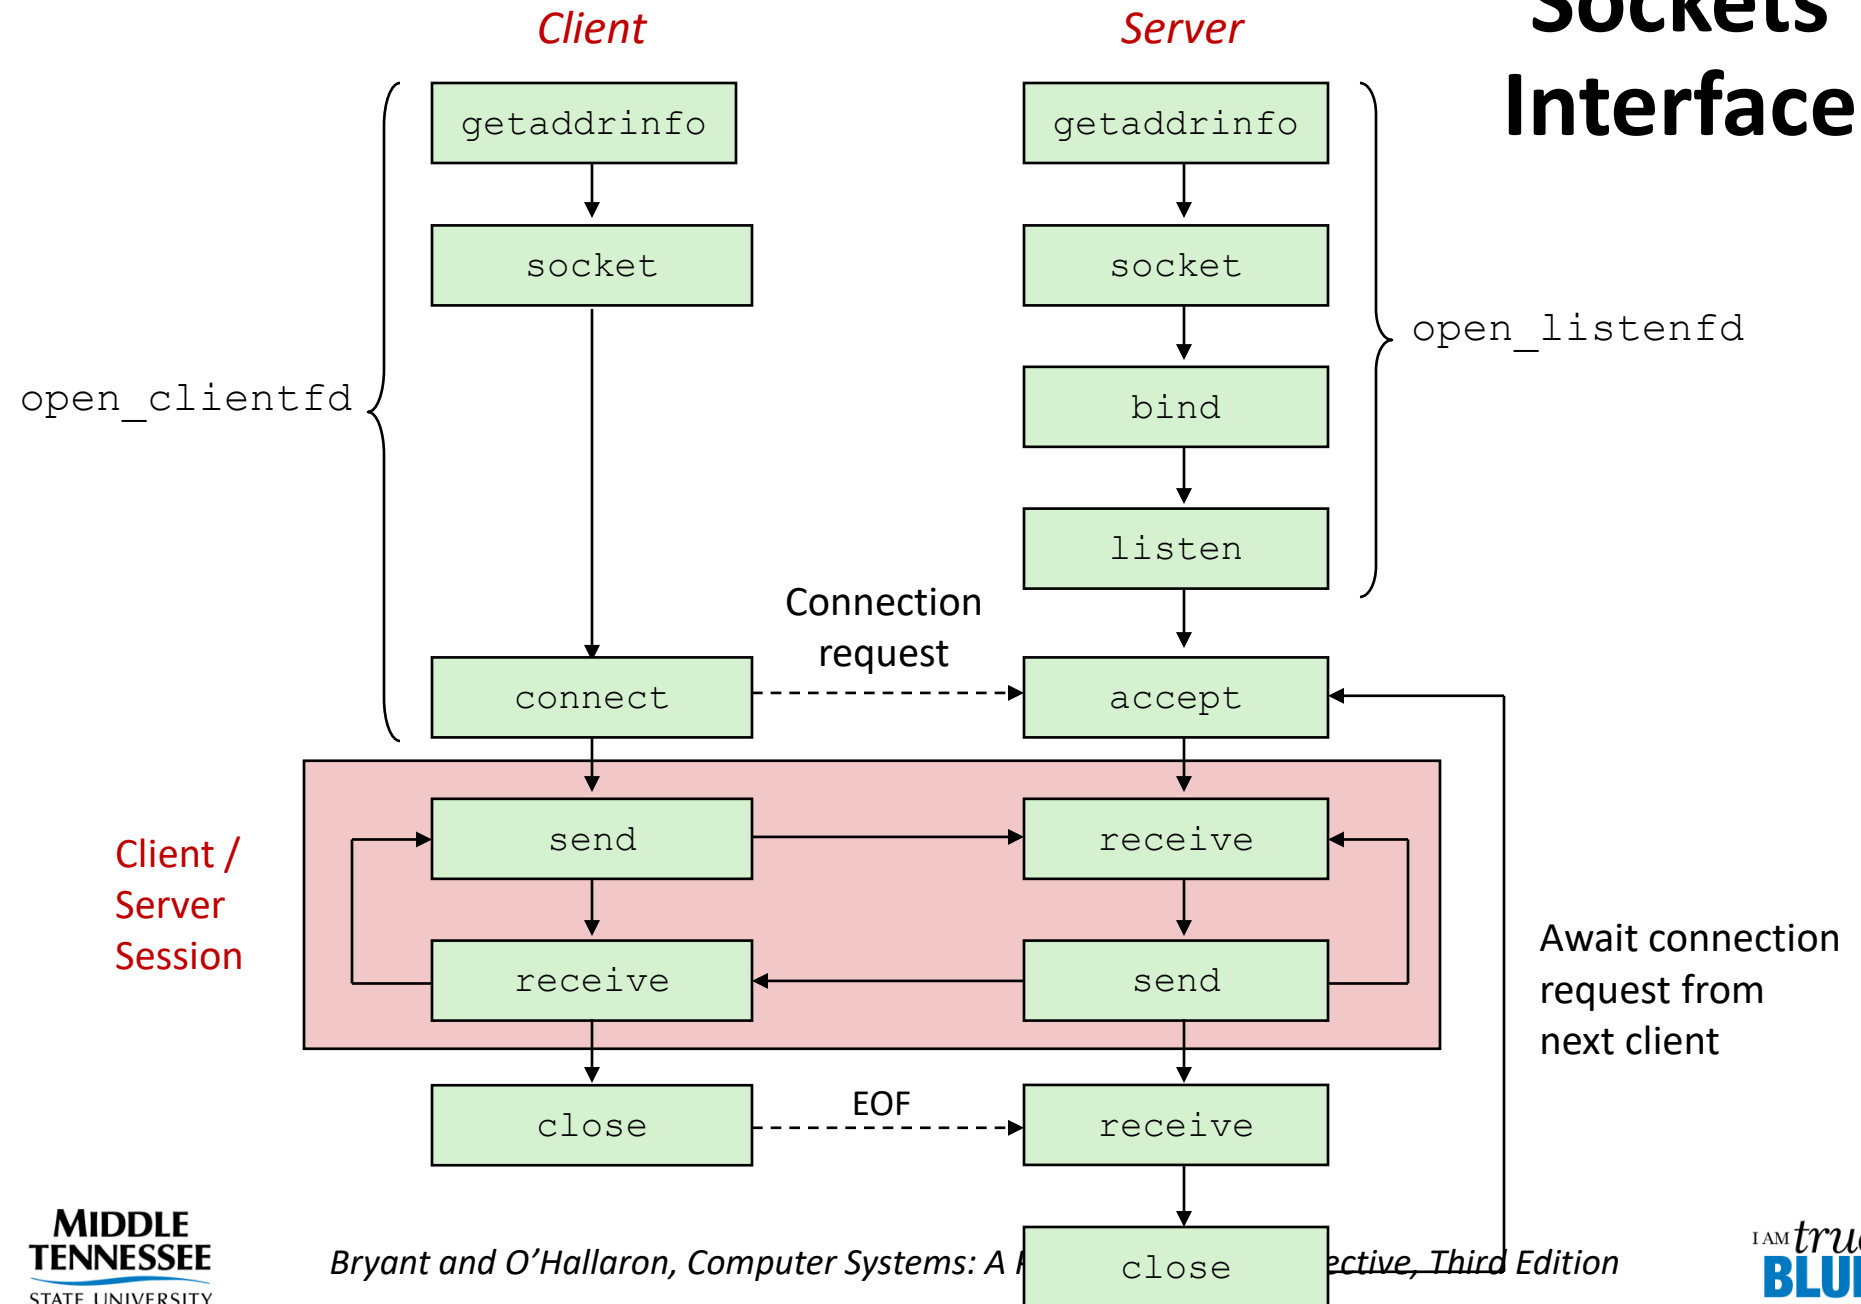
■ Connected descriptor

- End point of the connection between client and server
- A new descriptor is created each time the server accepts a connection request from a client
- Exists only as long as it takes to service client

■ Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously
 - E.g., Each time we receive a new request, we fork a child to handle the request

Sockets Interface



Host and Service Conversion: `getaddrinfo`

- `getaddrinfo` is the modern way to convert string representations of hostnames, host addresses, ports, and service names to socket address structures.
 - Replaces obsolete `gethostbyname` and `getservbyname` funcs.
- **Advantages:**
 - Reentrant (can be safely used by threaded programs).
 - Allows us to write portable protocol-independent code
 - Works with both IPv4 and IPv6
- **Disadvantages**
 - Somewhat complex
 - Fortunately, a small number of usage patterns suffice in most cases.

Host and Service Conversion: `getaddrinfo`

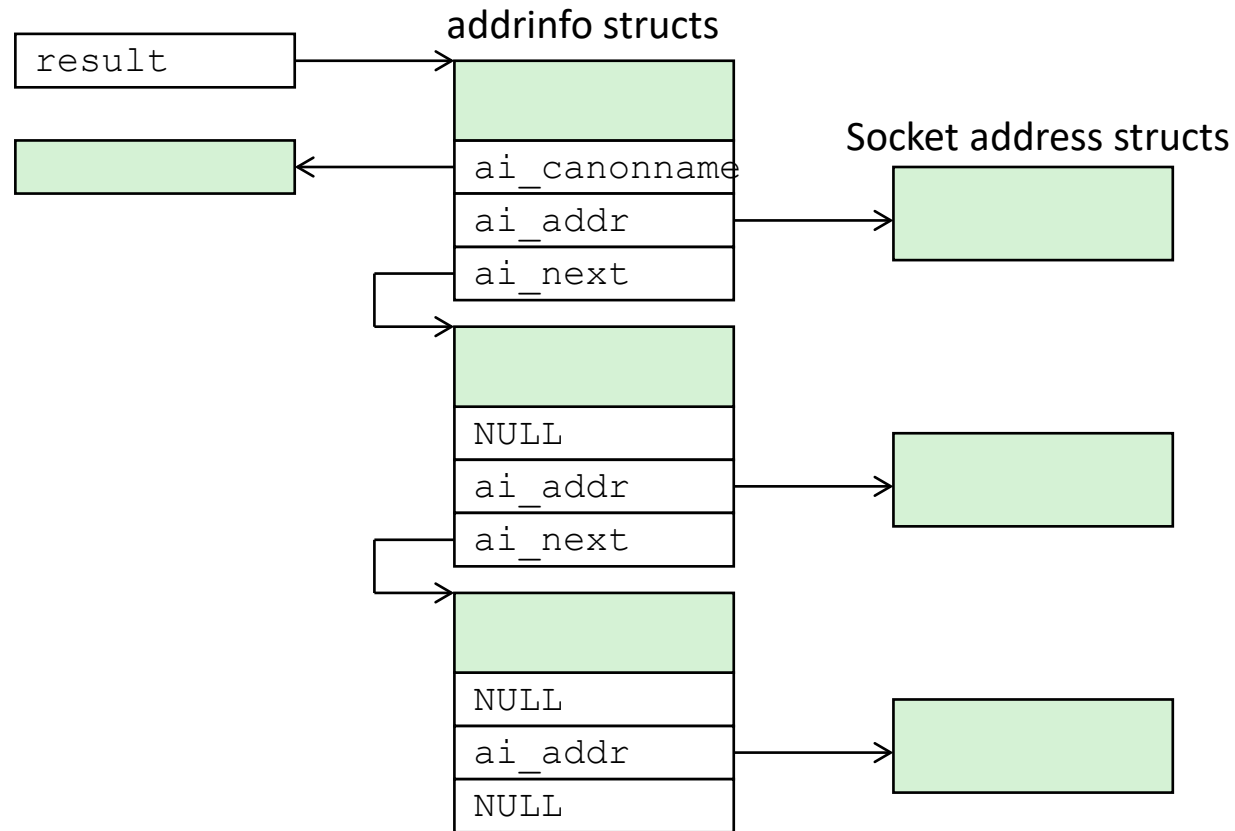
```
int getaddrinfo(const char *host,          /* Hostname or address */
                const char *service,      /* Port or service name */
                const struct addrinfo *hints, /* Input parameters */
                struct addrinfo **result); /* Output linked list */

void freeaddrinfo(struct addrinfo *result); /* Free linked list */

const char *gai_strerror(int errcode);    /* Return error msg */
```

- Given host and service, `getaddrinfo` returns result that points to a linked list of `addrinfo` structs, each of which points to a corresponding socket address struct, and which contains arguments for the sockets interface functions.
- **Helper functions:**
 - `freeaddrinfo` frees the entire linked list.
 - `gai_strerror` converts error code to an error message.

Linked List Returned by `getaddrinfo`



- **Clients:** walk this list, trying each socket address in turn, until the calls to `socket` and `connect` succeed.
- **Servers:** walk the list until calls to `socket` and `bind` succeed.

addrinfo Struct

```
struct addrinfo {  
    int             ai_flags;        /* Hints argument flags */  
    int             ai_family;       /* First arg to socket function */  
    int             ai_socktype;     /* Second arg to socket function */  
    int             ai_protocol;     /* Third arg to socket function */  
    char            *ai_canonname;    /* Canonical host name */  
    size_t          ai_addrlen;      /* Size of ai_addr struct */  
    struct sockaddr *ai_addr;         /* Ptr to socket address structure */  
    struct addrinfo *ai_next;         /* Ptr to next item in linked list */  
};
```

- Each **addrinfo** struct returned by **getaddrinfo** contains arguments that can be passed directly to **socket** function.
- Also points to a socket address struct that can be passed directly to **connect** and **bind** functions.

Host and Service Conversion: `getnameinfo`

- **`getnameinfo` is the inverse of `getaddrinfo`, converting a socket address to the corresponding host and service.**
 - Replaces obsolete `gethostbyaddr` and `getservbyport` funcs.
 - Reentrant and protocol independent.

```
int getnameinfo(const SA *sa, socklen_t salen, /* In: socket addr */
                char *host, size_t hostlen, /* Out: host */
                char *serv, size_t servlen, /* Out: service */
                int flags); /* optional flags */
```

Conversion Example

hostinfo.c

```
#include "csapp.h"

int main(int argc, char **argv)
{
    struct addrinfo *p, *listp, hints;
    char buf[MAXLINE];
    int rc, flags;

    /* Get a list of addrinfo records */
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;          /* IPv4 only */
    hints.ai_socktype = SOCK_STREAM; /* Connections only */
    if ((rc = getaddrinfo(argv[1], NULL, &hints, &listp)) != 0) {
        fprintf(stderr, "getaddrinfo error: %s\n", gai_strerror(rc));
        exit(1);
    }
}
```

hostinfo.c

Conversion Example (cont)

```
/* Walk the list and display each IP address */
flags = NI_NUMERICHOST; /* Display address instead of name */
for (p = listp; p; p = p->ai_next) {
    Getnameinfo(p->ai_addr, p->ai_addrlen,
                buf, MAXLINE, NULL, 0, flags);
    printf("%s\n", buf);
}

/* Clean up */
Freeaddrinfo(listp);

exit(0);
}
```

hostinfo.c

Running hostinfo

```
csci3240-00> ./hostinfo localhost
```

```
127.0.0.1
```

```
csci3240-00 > ./hostinfo www.cs.mtsu.edu
```

```
161.145.162.100
```

```
csci3240-00 > ./hostinfo google.com
```

```
142.250.10.147
```

```
142.250.10.104
```

```
142.250.10.106
```

```
142.250.10.99
```

```
142.250.10.103
```

```
142.250.10.105
```